

Chapter 7

Univariate Volatility Modeling

Alternative references for volatility modeling include chapters 10 and 11 in Taylor (2005), chapter 21 of Hamilton (1994), and chapter 4 of Enders (2004). Many of the original articles have been collected in Engle (1995).

Engle (1982) introduced the ARCH model and, in doing so, modern financial econometrics. Measuring and modeling conditional volatility is the cornerstone of the field. Models used for analyzing conditional volatility can be extended to capture a variety of related phenomena including Value-at-Risk, Expected Shortfall, forecasting the complete density of financial returns and duration analysis. This chapter begins by examining the meaning of “volatility” - it has many - before turning attention to the ARCH-family of models. The chapter details estimation, inference, model selection, forecasting, and diagnostic testing. The chapter concludes by covering new methods of measuring volatility: *realized volatility*, which makes use of using ultra-high-frequency data, and *implied volatility*, a measure of volatility computed from options prices.

Volatility measurement and modeling is the foundation of financial econometrics. This chapter begins by introducing volatility as a meaningful concept and then describes a widely used framework for volatility analysis: the ARCH model. The chapter describes the most widely used members of the ARCH family, fundamental properties of each, estimation, inference and model selection. Attention then turns to a new tool in the measurement and modeling of financial volatility, *realized volatility*, before concluding with a discussion of option-based *implied volatility*.

7.1 Why does volatility change?

Time-varying volatility is a pervasive empirical regularity in financial time series, and it is difficult to find an asset return series which does *not* exhibit time-varying volatility. This chapter focuses on providing a statistical description of the time-variation of volatility but does not go into depth on the economic causes of time-varying volatility. Many explanations have been proffered to explain this phenomenon, and treated individually; none provide a complete characterization of the variation in volatility observed in financial returns.

- *News Announcements*: The arrival of unanticipated news (or “news surprises”) forces agents to update beliefs. These new beliefs lead to portfolio rebalancing and high volatility correspond to periods when agents are incorporating the news and dynamically solving for new asset prices. While certain classes of assets have been shown to react to surprises, in particular, government bonds and foreign exchange, many appear to be unaffected by large surprises (see, *inter alia* Engle and Li (1998) and Andersen, Bollerslev, Diebold, and Vega (2007)). Additionally, news-induced periods of high volatility are generally short, often on the magnitude of 5 to 30-minutes and the apparent resolution of uncertainty is far too quick to explain the time-variation of volatility seen in asset prices.
- *Leverage*: When a firm is financed using both debt and equity, only the equity reflects the volatility of the firm’s cash flows. However, as the price of equity falls, the reduced equity must reflect the same volatility of the firm’s cash flows and so negative returns should lead to increases in equity volatility. The leverage effect is pervasive in equity returns, especially in broad equity indices, although alone it is insufficient to explain the time variation of volatility (Christie, 1982; Bekaert and Wu, 2000).
- *Volatility Feedback*: Volatility feedback is motivated by a model where the volatility of an asset is priced. When the price of an asset falls, the volatility must increase to reflect the increased expected return (in the future) of this asset, and an increase in volatility requires an even lower price to generate a sufficient return to compensate an investor for holding a volatile asset. There is evidence that empirically supports this explanation although this feature alone cannot explain the totality of the time-variation of volatility (Bekaert and Wu, 2000).
- *Illiquidity*: Short run spells of illiquidity may produce time-varying volatility even when shocks are i.i.d. Intuitively, if the market is oversold (bought), a small negative (positive) shock produces a small decrease (increase) in demand. However, since few participants are willing to buy (sell), this shock has a disproportionate effect on prices. Liquidity runs tend to last from 20 minutes to a few days and cannot explain the long cycles in present volatility.
- *State Uncertainty*: Asset prices are essential instruments that allow agents to express beliefs about the state of the economy. When the state is uncertain, slight changes in beliefs may cause significant shifts in portfolio holdings which in turn feedback into beliefs about the state. This feedback loop can generate time-varying volatility and should have the most substantial effect when the economy is transitioning between periods of growth and contraction (Veronesi, 1999; Collard et al., 2018).

The economic causes of the time-variation in volatility include all of these and some not yet identified, such as behavioral causes.

7.1.1 What is volatility?

Volatility comes in many shapes and forms. It is critical to distinguish between related but different uses of “volatility”.

Volatility Volatility is the standard deviation. Volatility is often preferred to variance as it is measured in the same *units* as the original data. For example, when using returns, the volatility is also measured in returns, and so volatility of 5% indicates that $\pm 5\%$ is a meaningful quantity.

Realized Volatility Realized volatility has historically been used to denote a measure of the volatility over some arbitrary period of time,

$$\hat{\sigma} = \sqrt{T^{-1} \sum_{t=1}^T (r_t - \hat{\mu})^2} \quad (7.1)$$

but is now used to describe a volatility measure constructed using ultra-high-frequency (UHF) data (also known as tick data). See section 7.8 for details.

Conditional Volatility Conditional volatility is the expected volatility at some future time $t + h$ based on all available information up to time t (\mathcal{F}_t). The one-period ahead conditional volatility is denoted $E_t[\sigma_{t+1}]$.

Implied Volatility Implied volatility is the volatility that correctly prices an option. The Black-Scholes pricing formula relates the price of a European call option to the current price of the underlying, the strike, the risk-free rate, the time-to-maturity, and the *volatility*,

$$BS(S_t, K, r, t, \sigma_t) = C_t$$

where C is the price of the call. The implied volatility is the value which solves the Black-Scholes taking the option and underlying prices, the strike, the risk-free and the time-to-maturity as given,

$$\hat{\sigma}_t(S_t, K, r, t, C).$$

Recent econometric developments have produced nonparametric estimators that do not make strong assumptions on the underlying price process. The VIX is a leading example of these these Model-free Implied Volatility (MFIV) estimators.

Annualized Volatility When volatility is measured over an interval other than a year, such as a day, week or month, it can always be scaled to reflect the volatility of the asset over a year. For example, if σ denotes the daily volatility of an asset and there are 252 trading days in a year, the annualized volatility is $\sqrt{252}\sigma$. Annualized volatility is a useful measure that removes the sampling interval from reported volatilities.

Variance All of the uses of volatility can be replaced with variance, and this chapter focuses on modeling *conditional variance* denoted $E_t[\sigma_{t+1}^2]$, or in the general case, $E_t[\sigma_{t+h}^2]$.

7.2 ARCH Models

In financial econometrics, an arch is not an architectural feature of a building; it is a fundamental tool for analyzing the time-variation of conditional variance. The success of the *ARCH* (AutoRegressive Conditional Heteroskedasticity) family of models can be attributed to three features: ARCH processes are essentially ARMA models, and many of the tools of linear time-series analysis can be directly applied, ARCH-family models are easy to estimate, and simple, parsimonious models are capable of accurate descriptions of the dynamics of asset volatility.

7.2.1 The ARCH model

The complete ARCH(P) model (Engle, 1982) relates the current level of volatility to the past P squared shocks.

Definition 7.1 (Pth Order Autoregressive Conditional Heteroskedasticity (ARCH)). A Pth order ARCH process is given by

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \alpha_2 \epsilon_{t-2}^2 + \dots + \alpha_P \epsilon_{t-P}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1). \end{aligned} \tag{7.2}$$

where μ_t can be any adapted model for the conditional mean.¹

The key feature of this model is that the variance of the shock, ϵ_t , is time varying and depends on the past P shocks, $\epsilon_{t-1}, \epsilon_{t-2}, \dots, \epsilon_{t-P}$, through their squares. σ_t^2 is the time $t - 1$ conditional variance. All of the right-hand side variables that determine σ_t^2 are known at time $t - 1$, and so σ_t^2 is in the time $t - 1$ information set \mathcal{F}_{t-1} . The model for the conditional mean can include own lags, shocks (in an MA model) or exogenous variables such as the default spread or term premium. In practice, the model for the conditional mean should be general enough to capture the dynamics present in the data. In many financial time series, particularly when returns are measured over short intervals - one day to one week - a constant mean, sometimes assumed to be 0, is sufficient.

An common alternative description an ARCH(P) model is

$$\begin{aligned} r_t | \mathcal{F}_{t-1} &\sim N(\mu_t, \sigma_t^2) \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \alpha_2 \epsilon_{t-2}^2 + \dots + \alpha_P \epsilon_{t-P}^2 \\ \epsilon_t &= r_t - \mu_t \end{aligned} \tag{7.3}$$

which is read “ r_t given the information set at time $t - 1$ is conditionally normal with mean μ_t and variance σ_t^2 ”.²

The conditional variance, σ_t^2 , is

$$\mathbb{E}_{t-1} [\epsilon_t^2] = \mathbb{E}_{t-1} [e_t^2 \sigma_t^2] = \sigma_t^2 \mathbb{E}_{t-1} [e_t^2] = \sigma_t^2 \tag{7.4}$$

and the *unconditional* variance, $\bar{\sigma}^2$, is

$$\mathbb{E} [\epsilon_{t+1}^2] = \bar{\sigma}^2. \tag{7.5}$$

¹A model is adapted if everything required to model the mean at time t is known at time $t - 1$. Standard examples of adapted mean processes include a constant mean, ARMA processes or models containing exogenous regressors known at time $t - 1$.

²It is implausible that the unconditional (long-run) mean return of many risky assets is zero. However, when using daily equity data, the squared mean is typically less than 1% of the variance ($\frac{\mu^2}{\sigma^2} < 0.01$) and there are few consequences for setting the conditional mean to 0. Some assets, e.g., electricity prices, have non-trivial predictability and an appropriate model for the conditional mean is required before modeling the volatility.

The first interesting property of the ARCH(P) model is the unconditional variance. Assuming the unconditional variance exists, $\bar{\sigma}^2 = E[\sigma_t^2]$ can be derived from

$$\begin{aligned}
 E[\sigma_t^2] &= E[\omega + \alpha_1 \epsilon_{t-1}^2 + \alpha_2 \epsilon_{t-2}^2 + \dots + \alpha_p \epsilon_{t-p}^2] \\
 &= \omega + \alpha_1 E[\epsilon_{t-1}^2] + \alpha_2 E[\epsilon_{t-2}^2] + \dots + \alpha_p E[\epsilon_{t-p}^2] \\
 &= \omega + \alpha_1 E[E_{t-2}[\sigma_{t-1}^2 e_{t-1}^2]] + \alpha_2 E[E_{t-3}[\sigma_{t-2}^2 e_{t-2}^2]] \\
 &\quad + \dots + \alpha_p E[E_{t-p-1}[\sigma_{t-p}^2 e_{t-p}^2]] \\
 &= \omega + \alpha_1 E[\sigma_{t-1}^2 E_{t-2}[e_{t-1}^2]] + \alpha_2 E[\sigma_{t-2}^2 E_{t-3}[e_{t-2}^2]] \\
 &\quad + \dots + \alpha_p E[\sigma_{t-p}^2 E_{t-p-1}[e_{t-p}^2]] \\
 &= \omega + \alpha_1 E[\sigma_{t-1}^2 \times 1] + \alpha_2 E[\sigma_{t-2}^2 \times 1] + \dots + \alpha_p E[\sigma_{t-p}^2 \times 1] \\
 &= \omega + \alpha_1 E[\sigma_{t-1}^2] + \alpha_2 E[\sigma_{t-2}^2] + \dots + \alpha_p E[\sigma_{t-p}^2] \\
 &= \omega + \alpha_1 E[\sigma_t^2] + \alpha_2 E[\sigma_t^2] + \dots + \alpha_p E[\sigma_t^2]
 \end{aligned} \tag{7.6}$$

$$E[\sigma_t^2] (1 - \alpha_1 - \alpha_2 - \dots - \alpha_p) = \omega$$

$$\bar{\sigma}^2 = \frac{\omega}{1 - \alpha_1 - \alpha_2 - \dots - \alpha_p}. \tag{7.7}$$

This derivation makes use of a number of properties of ARCH family models. First, the definition of the shock $\epsilon_t^2 \equiv e_t^2 \sigma_t^2$ is used to separate the i.i.d. normal innovation (e_t) from the conditional variance (σ_t^2) using the Law of Iterated Expectations. For example, σ_{t-1} is known at time $t-2$ and so is in \mathcal{F}_{t-2} and $E_{t-1}[\sigma_{t-1}] = \sigma_{t-1}$. e_{t-1} is an i.i.d. draw at time $t-1$, a random variable at time $t-2$, and so $E_{t-2}[e_{t-1}^2] = 1$. The result follows from the property that the unconditional expectation of σ_{t-j}^2 is the same in any time period ($E[\sigma_t^2] = E[\sigma_{t-p}^2] = \bar{\sigma}^2$) in a covariance stationary time series. Inspection of the final line in the derivation reveals the condition needed to ensure that the unconditional expectation is finite: $1 - \alpha_1 - \alpha_2 - \dots - \alpha_p > 0$. As was the case in an AR model, as the persistence (as measured by $\alpha_1, \alpha_2, \dots$) increases towards a unit root, the process explodes.

7.2.1.1 Stationarity

An ARCH(P) model is covariance stationary as long as the model for the conditional mean corresponds to a stationary process³ and $1 - \alpha_1 - \alpha_2 - \dots - \alpha_p > 0$.⁴ ARCH models have the property that $E[\epsilon_t^2] = \bar{\sigma}^2 = \omega / (1 - \alpha_1 - \alpha_2 - \dots - \alpha_p)$ since

$$E[\epsilon_t^2] = E[e_t^2 \sigma_t^2] = E[E_{t-1}[e_t^2 \sigma_t^2]] = E[\sigma_t^2 E_{t-1}[e_t^2]] = E[\sigma_t^2 \times 1] = E[\sigma_t^2]. \tag{7.8}$$

which exploits the conditional (on \mathcal{F}_{t-1}) independence of e_t from σ_t^2 and the assumption that e_t is a mean zero process with unit variance so that $E[e_t^2] = 1$.

One crucial requirement of any covariance stationary ARCH process is that the parameters of the variance evolution, $\alpha_1, \alpha_2, \dots, \alpha_p$ must all be positive.⁵ The intuition behind this requirement

³For example, a constant or a covariance stationary ARMA process.

⁴When $\sum_{i=1}^p \alpha_i > 1$, and ARCH(P) may still be strictly stationary although it cannot be covariance stationary since it has infinite variance.

⁵Since each $\alpha_j \geq 0$, the roots of the characteristic polynomial associated with $\alpha_1, \alpha_2, \dots, \alpha_p$ are less than 1 if and only if $\sum_{p=1}^p \alpha_p < 1$.

is that if one of the α s were negative, eventually a shock would be sufficiently large to produce a negative conditional variance and an ill-defined process. Finally, it is also necessary that $\omega > 0$ to ensure covariance stationarity.

To aid in developing intuition about ARCH-family models consider a simple ARCH(1) with a constant mean of 0,

$$\begin{aligned} r_t &= \epsilon_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1). \end{aligned} \tag{7.9}$$

While the conditional variance of an ARCH process appears different from anything previously encountered, the squared error ϵ_t^2 can be equivalently expressed as an AR(1). This transformation allows many properties of ARCH residuals to be directly derived by applying the results of chapter 4. By adding $\epsilon_t^2 - \sigma_t^2$ to both sides of the volatility equation,

$$\begin{aligned} \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 \\ \sigma_t^2 + \epsilon_t^2 - \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \epsilon_t^2 - \sigma_t^2 \\ \epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \epsilon_t^2 - \sigma_t^2 \\ \epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \sigma_t^2 (e_t^2 - 1) \\ \epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \nu_t, \end{aligned} \tag{7.10}$$

an ARCH(1) process can be shown to be an AR(1). The error term, ν_t represents the volatility *surprise*, $\epsilon_t^2 - \sigma_t^2$, which can be decomposed as $\sigma_t^2(e_t^2 - 1)$. The shock is a mean 0 white noise process since e_t is i.i.d. and $E[e_t^2] = 1$. Using the AR representation, the autocovariances of ϵ_t^2 are simple to derive. First note that $\epsilon_t^2 - \bar{\sigma}^2 = \sum_{i=0}^{\infty} \alpha_1^i \nu_{t-i}$. The first autocovariance can be expressed

$$\begin{aligned} E[(\epsilon_t^2 - \bar{\sigma}^2)(\epsilon_{t-1}^2 - \bar{\sigma}^2)] &= E\left[\left(\sum_{i=0}^{\infty} \alpha_1^i \nu_{t-i}\right)\left(\sum_{j=1}^{\infty} \alpha_1^{j-1} \nu_{t-j}\right)\right] \\ &= E\left[\left(\nu_t + \sum_{i=1}^{\infty} \alpha_1^i \nu_{t-i}\right)\left(\sum_{j=1}^{\infty} \alpha_1^{j-1} \nu_{t-j}\right)\right] \\ &= E\left[\left(\nu_t + \alpha_1 \sum_{i=1}^{\infty} \alpha_1^{i-1} \nu_{t-i}\right)\left(\sum_{j=1}^{\infty} \alpha_1^{j-1} \nu_{t-j}\right)\right] \\ &= E\left[\nu_t \left(\sum_{i=1}^{\infty} \alpha_1^{i-1} \nu_{t-i}\right)\right] + E\left[\alpha_1 \left(\sum_{i=1}^{\infty} \alpha_1^{i-1} \nu_{t-i}\right)\left(\sum_{j=1}^{\infty} \alpha_1^{j-1} \nu_{t-j}\right)\right] \\ &= \sum_{i=1}^{\infty} \alpha_1^{i-1} E[\nu_t \nu_{t-i}] + E\left[\alpha_1 \left(\sum_{i=1}^{\infty} \alpha_1^{i-1} \nu_{t-i}\right)\left(\sum_{j=1}^{\infty} \alpha_1^{j-1} \nu_{t-j}\right)\right] \end{aligned} \tag{7.11}$$

$$\begin{aligned}
&= \sum_{i=1}^{\infty} \alpha_1^{i-1} \cdot 0 + E \left[\alpha_1 \left(\sum_{i=1}^{\infty} \alpha_1^{i-1} v_{t-i} \right) \left(\sum_{j=1}^{\infty} \alpha_1^{j-1} v_{t-j} \right) \right] \\
&= \alpha_1 E \left[\left(\sum_{i=1}^{\infty} \alpha_1^{i-1} v_{t-i} \right)^2 \right] \\
&= \alpha_1 E \left[\left(\sum_{i=0}^{\infty} \alpha_1^i v_{t-1-i} \right)^2 \right] \\
&= \alpha_1 \left(\sum_{i=0}^{\infty} \alpha_1^{2i} E[v_{t-1-i}^2] + 2 \sum_{j=0}^{\infty} \sum_{k=j+1}^{\infty} \alpha_1^{jk} E[v_{t-1-j} v_{t-1-k}] \right) \\
&= \alpha_1 \left(\sum_{i=0}^{\infty} \alpha_1^{2i} V[v_{t-1-i}] + 2 \sum_{j=0}^{\infty} \sum_{k=j+1}^{\infty} \alpha_1^{jk} \cdot 0 \right) \\
&= \alpha_1 \sum_{i=0}^{\infty} \alpha_1^{2i} V[v_{t-1-i}] \\
&= \alpha_1 V[\epsilon_{t-1}^2]
\end{aligned}$$

where $V[\epsilon_{t-1}^2] = V[\epsilon_t^2]$ is the variance of the squared innovations.⁶ By repeated substitution, the s^{th} autocovariance, $E[(\epsilon_t^2 - \bar{\sigma}^2)(\epsilon_{t-s}^2 - \bar{\sigma}^2)]$, can be shown to be $\alpha_1^s V[\epsilon_t^2]$, and so that the autocovariances of an ARCH(1) process are identical to those of an AR(1) process.

7.2.1.2 Autocorrelations

Using the autocovariances, the autocorrelations are

$$\text{Corr}(\epsilon_t^2, \epsilon_{t-s}^2) = \frac{\alpha_1^s V[\epsilon_t^2]}{V[\epsilon_t^2]} = \alpha_1^s. \quad (7.12)$$

Further, the relationship between the s^{th} autocorrelation of an ARCH process and an AR process holds for ARCH processes with other orders. The autocorrelations of an ARCH(P) are identical to those of an AR(P) process with $\{\phi_1, \phi_2, \dots, \phi_P\} = \{\alpha_1, \alpha_2, \dots, \alpha_P\}$. One interesting aspect of ARCH(P) processes (and any covariance stationary ARCH-family model) is that the autocorrelations of $\{\epsilon_t^2\}$ *must* be positive. If one autocorrelation were negative, eventually a shock would be sufficiently large to force the conditional variance negative, and so the process would be ill-defined. In practice it is often better to examine the absolute values ($\text{Corr}(|\epsilon_t|, |\epsilon_{t-s}|)$) rather than the squares since financial returns frequently have outliers that are exacerbated when squared.

7.2.1.3 Kurtosis

The second interesting property of ARCH models is that the kurtosis of shocks (ϵ_t) is strictly greater than the kurtosis of a normal. This may seem strange since all of the shocks $\epsilon_t = \sigma_t e_t$

⁶For the time being, assume this is finite.




Figure 7.1: Plots of S&P 500 and WTI returns (scaled by 100) from 1999 until 2018. The bulges in the return plots are graphical evidence of time-varying volatility.

are normal by assumption. An ARCH model is a *variance-mixture* of normals, and so must have a kurtosis larger than three. The direct proof is simple,

$$\kappa = \frac{E[\epsilon_t^4]}{E[\epsilon_t^2]^2} = \frac{E[E_{t-1}[\epsilon_t^4]]}{E[E_{t-1}[e_t^2 \sigma_t^2]]^2} = \frac{E[E_{t-1}[e_t^4 \sigma_t^4]]}{E[E_{t-1}[e_t^2 \sigma_t^2]]^2} = \frac{E[3\sigma_t^4]}{E[\sigma_t^2]^2} = 3 \frac{E[\sigma_t^4]}{E[\sigma_t^2]^2} \geq 3. \quad (7.13)$$

The key steps in this derivation are that $\epsilon_t^4 = e_t^4 \sigma_t^4$ and that $E_t[e_t^4] = 3$ since e_t is a standard normal. The final conclusion that $E[\sigma_t^4]/E[\sigma_t^2]^2 > 1$ follows from noting that for any random variable Y , $V[Y] = E[Y^2] - E[Y]^2 \geq 0$ and so it must be the case that $E[\sigma_t^4] \geq E[\sigma_t^2]^2$ or $\frac{E[\sigma_t^4]}{E[\sigma_t^2]^2} \geq 1$. The kurtosis, κ , of an ARCH(1) can be shown to be

$$\kappa = \frac{3(1 - \alpha_1^2)}{(1 - 3\alpha_1^2)} > 3 \quad (7.14)$$

which is greater than 3 since $1 - 3\alpha_1^2 < 1 - \alpha_1^2$ for any value of $\alpha \neq 0$. The complete derivation of the kurtosis is involved and is presented in Appendix 7.A.




Figure 7.2: Plots of the squared returns of the S&P 500 Index and WTI. Time-variation in the squared returns is evidence of ARCH.

7.2.2 The GARCH model

The ARCH model has been deemed a sufficient contribution to economics to warrant a Nobel prize. Unfortunately, like most models, it has problems. ARCH models typically require 5-8 lags of the squared shock to model conditional variance adequately. The Generalized ARCH (GARCH) process, introduced by Bollerslev (1986), improves the original specification adding lagged conditional variance, which acts as a *smoothing* term. A low-order GARCH model typically fits as well as a high-order ARCH.

Definition 7.2 (Generalized Autoregressive Conditional Heteroskedasticity (GARCH) process). A GARCH(P,Q) process is defined as

$$\begin{aligned}
 r_t &= \mu_t + \epsilon_t \\
 \sigma_t^2 &= \omega + \sum_{p=1}^P \alpha_p \epsilon_{t-p}^2 + \sum_{q=1}^Q \beta_q \sigma_{t-q}^2 \\
 \epsilon_t &= \sigma_t e_t
 \end{aligned} \tag{7.15}$$

$$e_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$$

where μ_t can be any adapted model for the conditional mean.

The GARCH(P,Q) model builds on the ARCH(P) model by including Q lags of the conditional variance, $\sigma_{t-1}^2, \sigma_{t-2}^2, \dots, \sigma_{t-Q}^2$. Rather than focusing on the general specification with all of its complications, consider a simpler GARCH(1,1) model where the conditional mean is assumed to be zero,

$$\begin{aligned} r_t &= \epsilon_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned} \tag{7.16}$$

In this specification, the future variance will be an average of the current shock, ϵ_{t-1}^2 , the current variance, σ_{t-1}^2 , and a constant. Including the lagged variance produces a model that can be equivalently expressed as an ARCH(∞). Begin by backward substituting for σ_{t-1}^2 ,

$$\begin{aligned} \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \\ &= \omega + \alpha_1 \epsilon_{t-1}^2 + \underbrace{\beta_1 (\omega + \alpha_1 \epsilon_{t-2}^2 + \beta_1 \sigma_{t-2}^2)}_{\sigma_{t-1}^2} \\ &= \omega + \beta_1 \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \alpha_1 \epsilon_{t-2}^2 + \beta_1^2 \sigma_{t-2}^2 \\ &= \omega + \beta_1 \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \alpha_1 \epsilon_{t-2}^2 + \underbrace{\beta_1^2 \omega + \alpha_1 \epsilon_{t-3}^2 + \beta_1 \sigma_{t-3}^2}_{\sigma_{t-2}^2} \\ &= \omega + \beta_1 \omega + \beta_1^2 \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \alpha_1 \epsilon_{t-2}^2 + \beta_1^2 \alpha_1 \epsilon_{t-3}^2 + \beta_1^3 \sigma_{t-3}^2 \\ &= \sum_{i=0}^{\infty} \beta_1^i \omega + \sum_{i=0}^{\infty} \beta_1^i \alpha_1 \epsilon_{t-i-1}^2, \end{aligned} \tag{7.17}$$

and the ARCH(∞) representation can be derived.⁷ The conditional variance in period t depends on a constant, $\sum_{i=0}^{\infty} \beta_1^i \omega = \frac{\omega}{1-\beta}$, and a weighted average of past squared innovations with weights $\alpha_1, \beta_1 \alpha_1, \beta_1^2 \alpha_1, \beta_1^3 \alpha_1, \dots$

As was the case in the ARCH(P) model, the coefficients of a GARCH model must be restricted to ensure the conditional variances are uniformly positive. In a GARCH(1,1), these restrictions are $\omega > 0$, $\alpha_1 \geq 0$ and $\beta_1 \geq 0$. In a GARCH(P,1) model the restriction change to $\alpha_p \geq 0$, $p = 1, 2, \dots, P$ with the same restrictions on ω and β_1 . The minimal parameter restrictions needed to ensure that variances are always positive are difficult to derive for the full class of GARCH(P,Q) models. For example, in a GARCH(2,2), one of the two β 's (β_2) can be slightly negative while ensuring that all conditional variances are positive. See Nelson and Cao (1992) for further details.

The GARCH(1,1) model can be transformed into a standard time series model for ϵ_t^2 by adding $\epsilon_t^2 - \sigma_t^2$ to both sides.

⁷Since the model is assumed to be stationary, it must be the case that $0 \leq \beta < 1$ and so $\lim_{j \rightarrow \infty} \beta^j \sigma_{t-j} = 0$.

$$\begin{aligned}
\sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \\
\sigma_t^2 + \epsilon_t^2 - \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 + \epsilon_t^2 - \sigma_t^2 \\
\epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 + \epsilon_t^2 - \sigma_t^2 \\
\epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 - \beta_1 \epsilon_{t-1}^2 + \beta_1 \epsilon_{t-1}^2 + \epsilon_t^2 - \sigma_t^2 \\
\epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \epsilon_{t-1}^2 - \beta_1 (\epsilon_{t-1}^2 - \sigma_{t-1}^2) + \epsilon_t^2 - \sigma_t^2 \\
\epsilon_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \epsilon_{t-1}^2 - \beta_1 \nu_{t-1} + \nu_t \\
\epsilon_t^2 &= \omega + (\alpha_1 + \beta_1) \epsilon_{t-1}^2 - \beta_1 \nu_{t-1} + \nu_t
\end{aligned} \tag{7.18}$$

The squared shock in a GARCH(1,1) follows an ARMA(1,1) process where $\nu_t = \epsilon_t^2 - \sigma_t^2$ is the volatility surprise. In the general GARCH(P,Q), the ARMA representation takes the form of an ARMA(max(P,Q),Q).

$$\epsilon_t^2 = \omega + \sum_{i=1}^{\max(P,Q)} (\alpha_i + \beta_i) \epsilon_{t-i}^2 - \sum_{q=1}^Q \beta_q \nu_{t-q} + \nu_t \tag{7.19}$$

The unconditional variance is computed by taking expectations of both sides, so that

$$\begin{aligned}
E[\sigma_t^2] &= \omega + \alpha_1 E[\epsilon_{t-1}^2] + \beta_1 E[\sigma_{t-1}^2] \\
\bar{\sigma}^2 &= \omega + \alpha_1 \bar{\sigma}^2 + \beta_1 \bar{\sigma}^2 \\
\bar{\sigma}^2 - \alpha_1 \bar{\sigma}^2 - \beta_1 \bar{\sigma}^2 &= \omega \\
\bar{\sigma}^2 &= \frac{\omega}{1 - \alpha_1 - \beta_1}.
\end{aligned} \tag{7.20}$$

Inspection of the ARMA model leads to an alternative derivation of $\bar{\sigma}^2$ since the AR coefficient is $\alpha_1 + \beta_1$ and the intercept is ω , and the unconditional mean in an ARMA(1,1) is the intercept divided by one minus the AR coefficient, $\omega/(1 - \alpha_1 - \beta_1)$. In a general GARCH(P,Q) the unconditional variance is

$$\bar{\sigma}^2 = \frac{\omega}{1 - \sum_{p=1}^P \alpha_p - \sum_{q=1}^Q \beta_q}. \tag{7.21}$$

The requirements on the parameters for stationarity in a GARCH(1,1) are $1 - \alpha_1 - \beta_1 > 0$ and $\alpha_1 \geq 0, \beta_1 \geq 0$ and $\omega > 0$.

The ARMA(1,1) form can be used directly to solve for the autocovariances. Recall the definition of a mean zero ARMA(1,1),

$$Y_t = \phi Y_{t-1} + \theta \epsilon_{t-1} + \epsilon_t \tag{7.22}$$

The 1st autocovariance can be computed as

$$\begin{aligned}
E[Y_t Y_{t-1}] &= E[(\phi Y_{t-1} + \theta \epsilon_{t-1} + \epsilon_t) Y_{t-1}] \\
&= E[\phi Y_{t-1}^2] + [\theta \epsilon_{t-1}^2] \\
&= \phi V[Y_{t-1}] + \theta V[\epsilon_{t-1}] \\
\gamma_1 &= \phi V[Y_{t-1}] + \theta V[\epsilon_{t-1}]
\end{aligned} \tag{7.23}$$

and the s^{th} autocovariance is $\gamma_s = \phi^{s-1}\gamma_1$. In the notation of a GARCH(1,1) model, $\phi = \alpha_1 + \beta_1$, $\theta = -\beta_1$, Y_{t-1} is ϵ_{t-1}^2 and η_{t-1} is $\sigma_{t-1}^2 - \epsilon_{t-1}^2$. Thus, $V[\epsilon_{t-1}^2]$ and $V[\sigma_t^2 - \epsilon_t^2]$ must be solved for. This derivation is challenging and so is presented in the appendix. The key to understanding the autocovariance (and autocorrelation) of a GARCH is to use the ARMA mapping. First note that $E[\sigma_t^2 - \epsilon_t^2] = 0$ so $V[\sigma_t^2 - \epsilon_t^2]$ is simply $E[(\sigma_t^2 - \epsilon_t^2)^2]$. This can be expanded to $E[\epsilon_t^4] - 2E[\epsilon_t^2 \sigma_t^2] + E[\sigma_t^4]$ which can be shown to be $2E[\sigma_t^4]$. The only remaining step is to complete the tedious derivation of the expectation of these fourth powers which is presented in Appendix 7.B.

7.2.2.1 Kurtosis

The kurtosis can be shown to be

$$\kappa = \frac{3(1 + \alpha_1 + \beta_1)(1 - \alpha_1 - \beta_1)}{1 - 2\alpha_1\beta_1 - 3\alpha_1^2 - \beta_1^2} > 3. \quad (7.24)$$

The kurtosis is larger than that of a normal despite the innovations, e_t , all having normal distributions since that model is a variance mixture of normals. The formal derivation is presented in 7.B.

Exponentially Weighted Moving Averages (EWMA)

Exponentially Weighted Moving Averages, popularized by RiskMetrics, are commonly used to measure and forecast volatilities from returns without estimating any parameters (J.P.Morgan/Reuters, 1996). An EWMA is a restricted GARCH(1,1) model where $\omega = 0$ and $\alpha + \beta = 1$. The recursive form of an EWMA is

$$\sigma_t^2 = (1 - \lambda)\epsilon_{t-1}^2 + \lambda\sigma_{t-1}^2,$$

which can be equivalently expressed as an ARCH(∞)

$$\sigma_t^2 = (1 - \lambda) \sum_{i=0}^{\infty} \lambda^i \epsilon_{t-i-1}^2.$$

The weights on the lagged squared returns decay exponentially so that the ratio of two consecutive weights is λ . The single parameter λ is typically set to 0.94 when using daily returns, 0.97 when using weekly return data, or 0.99 when using monthly returns. These values were calibrated on a wide range of assets to forecast volatility well.

7.2.3 The EGARCH model

The Exponential GARCH (EGARCH) model represents a major shift from the ARCH and GARCH models (Nelson, 1991). Rather than model the variance directly, EGARCH models the natural logarithm of the variance, and so no parameters restrictions are required to ensure that the conditional variance is positive.

Definition 7.3 (Exponential Generalized Autoregressive Conditional Heteroskedasticity (EGARCH))

process). An EGARCH(P,O,Q) process is defined

$$\begin{aligned}
 r_t &= \mu_t + \epsilon_t \\
 \ln(\sigma_t^2) &= \omega + \sum_{p=1}^P \alpha_p \left(\left| \frac{\epsilon_{t-p}}{\sigma_{t-p}} \right| - \sqrt{\frac{2}{\pi}} \right) + \sum_{o=1}^O \gamma_o \frac{\epsilon_{t-o}}{\sigma_{t-o}} + \sum_{q=1}^Q \beta_q \ln(\sigma_{t-q}^2) \\
 \epsilon_t &= \sigma_t e_t \\
 e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1)
 \end{aligned} \tag{7.25}$$

where μ_t can be any adapted model for the conditional mean. P and O were assumed to be equal in the original parameterization of Nelson (1991).

Rather than working with the complete specification, consider a simpler version, an EGARCH(1,1,1) with a constant mean,

$$\begin{aligned}
 r_t &= \mu + \epsilon_t \\
 \ln(\sigma_t^2) &= \omega + \alpha_1 \left(\left| \frac{\epsilon_{t-1}}{\sigma_{t-1}} \right| - \sqrt{\frac{2}{\pi}} \right) + \gamma_1 \frac{\epsilon_{t-1}}{\sigma_{t-1}} + \beta_1 \ln(\sigma_{t-1}^2) \\
 \epsilon_t &= \sigma_t e_t \\
 e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1).
 \end{aligned} \tag{7.26}$$

Three terms drive the dynamics in the log variance. The first term, $\left| \frac{\epsilon_{t-1}}{\sigma_{t-1}} \right| - \sqrt{\frac{2}{\pi}} = |e_{t-1}| - \sqrt{\frac{2}{\pi}}$, is the absolute value of a normal random variable, e_{t-1} , minus its expectation, $\sqrt{2/\pi}$, and so it is a mean zero shock. The second term, e_{t-1} – a standard normal – is an additional mean zero shock and the final term is the lagged log variance. The two shocks behave differently: the absolute value in the first produces a symmetric rise in the log variance for a given return while the sign of the second produces an asymmetric effect. γ_1 is typically estimated to be negative so that volatility rises more after negative shocks than after positive ones. In the usual case where $\gamma_1 < 0$, the magnitude of the shock can be decomposed by conditioning on the sign of e_{t-1}

$$\text{Shock coefficient} = \begin{cases} \alpha_1 + \gamma_1 & \text{when } e_{t-1} < 0 \\ \alpha_1 - \gamma_1 & \text{when } e_{t-1} > 0 \end{cases} \tag{7.27}$$

Since both shocks are mean zero and the current log variance is linearly related to past log variance through β_1 , the EGARCH(1,1,1) model is an AR model.

EGARCH models often provide superior fits when compared to standard GARCH models. The presence of the asymmetric term is largely responsible for the superior fit since many asset return series have been found to exhibit a “leverage” effect. Additionally, the use of standardized shocks (e_{t-1}) in the dynamics of the log-variance reduces the effect of outliers.

7.2.3.1 The S&P 500 and West Texas Intermediate Crude

The application of GARCH models will be demonstrated using daily returns on both the S&P 500 and West Texas Intermediate (WTI) Crude spot prices from January 1, 1999, until December 31,

Summary Statistics		
	S&P 500	WTI
Ann. Mean	14.03	5.65
Ann. Volatility	38.57	19.04
Skewness	0.063	-0.028
Kurtosis	7.22	11.45

Table 7.1: Summary statistics for the S&P 500 and WTI. Means and volatilities are reported in annualized terms using $100 \times$ returns. Skewness and kurtosis are scale-free by definition.

2018. The S&P 500 data is from Yahoo! Finance and the WTI data is from the St. Louis Federal Reserve's FRED database. All returns are scaled by 100. The returns are plotted in Figure 7.1, the squared returns are plotted in Figure 7.2, and the absolute values of the returns are plotted in Figure 7.3. The plots of the squared returns and the absolute values of the returns are useful graphical diagnostics for detecting ARCH. If the residuals are conditionally heteroskedastic, both plots provide evidence of volatility dynamics in the transformed returns. In practice, the plot of the absolute returns is a more helpful graphical tool than the plot of the squares. Squared returns are noisy proxies for the variance, and the dynamics in the data may be obscured by a small number of outliers.

Summary statistics are presented in table 7.1, and estimates from an ARCH(5), and GARCH(1,1) and an EGARCH(1,1,1) are presented in table 7.2. The summary statistics are typical of financial data where both series are heavy-tailed (leptokurtotic).

Definition 7.4 (Leptokurtosis). A random variable x_t is said to be leptokurtic if its kurtosis,

$$\kappa = \frac{E[(x_t - E[x_t])^4]}{E[(x_t - E[x_t])^2]^2}$$

is greater than that of a normal ($\kappa > 3$). Leptokurtic variables are also known as “heavy-tailed” or “fat tailed”.

Definition 7.5 (Platykurtosis). A random variable x_t is said to be platykurtic if its kurtosis,

$$\kappa = \frac{E[(x_t - E[x_t])^4]}{E[(x_t - E[x_t])^2]^2}$$

is less than that of a normal ($\kappa < 3$). Platykurtic variables are also known as “thin-tailed”.

Table 7.2 contains estimates from an ARCH(5), a GARCH(1,1) and an EGARCH(1,1,1) model. All estimates were computed using maximum likelihood assuming the innovations are conditionally normally distributed. There is strong evidence of time-varying variance since most p-values are near 0. The highest log-likelihood (a measure of fit) is produced by the EGARCH model in both series. This is likely due to the EGARCH's inclusion of asymmetries, a feature excluded from both the ARCH and GARCH models.

S&P 500						
ARCH(5)						
ω	α_1	α_2	α_3	α_4	α_5	Log Lik.
0.294 (0.000)	0.095 (0.000)	0.204 (0.000)	0.189 (0.000)	0.193 (0.000)	0.143 (0.000)	−7008
GARCH(1,1)						
ω	α_1	β_1				Log Lik.
0.018 (0.000)	0.102 (0.000)	0.885 (0.000)				−6888
EGARCH(1,1,1)						
ω	α_1	γ_1	β_1			Log Lik.
0.000 (0.909)	0.136 (0.000)	−0.153 (0.000)	0.975 (0.000)			−6767
WTI						
ARCH(5)						
ω	α_1	α_2	α_3	α_4	α_5	Log Lik.
2.282 (0.000)	0.138 (0.000)	0.129 (0.000)	0.131 (0.000)	0.094 (0.000)	0.130 (0.000)	−11129
GARCH(1,1)						
ω	α_1	β_1				Log Lik.
0.047 (0.034)	0.059 (0.000)	0.934 (0.000)				−11030
EGARCH(1,1,1)						
ω	α_1	γ_1	β_1			Log Lik.
0.020 (0.002)	0.109 (0.000)	−0.050 (0.000)	0.990 (0.000)			−11001

Table 7.2: Parameter estimates, p-values and log-likelihoods from ARCH(5), GARCH(1,1) and EGARCH(1,1,1) models for the S&P 500 and WTI. These parameter values are typical of models estimated on daily data. The persistence of conditional variance, as measured by the sum of the α s in the ARCH(5), $\alpha_1 + \beta_1$ in the GARCH(1,1) and β_1 in the EGARCH(1,1,1), is high in all models. The log-likelihoods indicate the EGARCH model is preferred for both return series.




Figure 7.3: Plots of the absolute returns of the S&P 500 and WTI. Plots of the absolute value are often more useful in detecting ARCH as they are less noisy than squared returns yet still show changes in conditional volatility.

7.2.4 Alternative Specifications

Many extensions to the basic ARCH model have been introduced to capture important empirical features. This section outlines three of the most useful extensions in the ARCH-family.

7.2.4.1 GJR-GARCH

The GJR-GARCH model was named after the authors who introduced it, Glosten, Jagannathan, and Runkle (1993). It extends the standard GARCH(P,Q) by adding asymmetric terms that capture a common phenomenon in the conditional variance of equities: the propensity of the volatility to rise more after large negative shocks than to large positive shocks (known as the “leverage effect”).

Definition 7.6 (GJR-Generalized Autoregressive Conditional Heteroskedasticity (GJR-GARCH))

process). A GJR-GARCH(P,O,Q) process is defined as

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t^2 &= \omega + \sum_{p=1}^P \alpha_p \epsilon_{t-p}^2 + \sum_{o=1}^O \gamma_o \epsilon_{t-o}^2 I_{[\epsilon_{t-o} < 0]} + \sum_{q=1}^Q \beta_q \sigma_{t-q}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned} \quad (7.28)$$

where μ_t can be any adapted model for the conditional mean, and $I_{[\epsilon_{t-o} < 0]}$ is an indicator function that takes the value 1 if $\epsilon_{t-o} < 0$ and 0 otherwise.

The parameters of the GJR-GARCH, like the standard GARCH model, must be restricted to ensure that the fit variances are always positive. This set is difficult to describe for all GJR-GARCH(P,O,Q) models although it is simple of a GJR-GARCH(1,1,1). The dynamics in a GJR-GARCH(1,1,1) evolve according to

$$\sigma_t^2 = \omega + \alpha_1 \epsilon_{t-1}^2 + \gamma_1 \epsilon_{t-1}^2 I_{[\epsilon_{t-1} < 0]} + \beta_1 \sigma_{t-1}^2. \quad (7.29)$$

and it must be the case that $\omega > 0$, $\alpha_1 \geq 0$, $\alpha_1 + \gamma \geq 0$ and $\beta_1 \geq 0$. If the innovations are conditionally normal, a GJR-GARCH model will be covariance stationary as long as the parameter restriction are satisfied and $\alpha_1 + \frac{1}{2}\gamma_1 + \beta_1 < 1$.

7.2.4.2 AVGARCH/TARCH/ZARCH

The Threshold ARCH (TARCH) model (also known as AVGARCH and ZARCH) makes one fundamental change to the GJR-GARCH model (Taylor, 1986; Zakoian, 1994). The TARCH model parameterizes the *conditional standard deviation* as a function of the lagged absolute value of the shocks. It also captures asymmetries using an asymmetric term that is similar to the asymmetry in the GJR-GARCH model.

Definition 7.7 (Threshold Autoregressive Conditional Heteroskedasticity (TARCH) process). A TARCH(P, O, Q) process is defined as

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t &= \omega + \sum_{p=1}^P \alpha_p |\epsilon_{t-p}| + \sum_{o=1}^O \gamma_o |\epsilon_{t-o}| I_{[\epsilon_{t-o} < 0]} + \sum_{q=1}^Q \beta_q \sigma_{t-q} \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned} \quad (7.30)$$

where μ_t can be any adapted model for the conditional mean. TARCH models are also known as ZARCH due to Zakoian (1994) or AVGARCH when no asymmetric terms are included ($O = 0$, Taylor (1986)).

Below is an example of a TARCH(1,1,1) model.

$$\sigma_t = \omega + \alpha_1 |\epsilon_{t-1}| + \gamma_1 |\epsilon_{t-1}| I_{[\epsilon_{t-1} < 0]} + \beta_1 \sigma_{t-1}, \quad \alpha_1 + \gamma_1 \geq 0 \quad (7.31)$$

where $I_{[\epsilon_{t-1} < 0]}$ is an indicator variable which takes the value 1 if $\epsilon_{t-1} < 0$. Models of the conditional standard deviation often outperform models that parameterize the conditional variance. The absolute shocks are less responsive than the squared shocks.

7.2.4.3 APARCH

The third model extends the TARCH and GJR-GARCH models by directly parameterizing the non-linearity in the conditional variance. Where the GJR-GARCH model uses 2, and the TARCH model uses 1, the Asymmetric Power ARCH (APARCH) of Ding, Granger, and Engle (1993) parameterizes this value directly (using δ). This form provides greater flexibility in modeling the memory of volatility while remaining parsimonious.

Definition 7.8 (Asymmetric Power Autoregressive Conditional Heteroskedasticity (APARCH) process). An APARCH(P,O,Q) process is defined as

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t^\delta &= \omega + \sum_{j=1}^{\max(P,O)} \alpha_j (|\epsilon_{t-j}| + \gamma_j \epsilon_{t-j})^\delta + \sum_{q=1}^Q \beta_q \sigma_{t-q}^\delta \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned} \quad (7.32)$$

where μ_t can be any adapted model for the conditional mean. It must be the case that $P \geq O$ in an APARCH model, and if $P > O$, then $\gamma_j = 0$ for $j > O$. If that $\omega > 0$, $\alpha_k \geq 0$ and $-1 \leq \gamma_j \leq 1$, then the conditional variance are always positive.

It is not completely obvious to see that the APARCH model nests the GJR-GARCH and TARCH models as special cases. To examine how an APARCH nests a GJR-GARCH, consider an APARCH(1,1,1) model.

$$\sigma_t^\delta = \omega + \alpha_1 (|\epsilon_{t-1}| + \gamma_1 \epsilon_{t-1})^\delta + \beta_1 \sigma_{t-1}^\delta \quad (7.33)$$

Suppose $\delta = 2$, then

$$\begin{aligned} \sigma_t^2 &= \omega + \alpha_1 (|\epsilon_{t-1}| + \gamma_1 \epsilon_{t-1})^2 + \beta_1 \sigma_{t-1}^2 \\ &= \omega + \alpha_1 |\epsilon_{t-1}|^2 + 2\alpha_1 \gamma_1 \epsilon_{t-1} |\epsilon_{t-1}| + \alpha_1 \gamma_1^2 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \\ &= \omega + \alpha_1 \epsilon_{t-1}^2 + \alpha_1 \gamma_1^2 \epsilon_{t-1}^2 + 2\alpha_1 \gamma_1 \epsilon_{t-1}^2 \text{sign}(\epsilon_{t-1}) + \beta_1 \sigma_{t-1}^2 \end{aligned} \quad (7.34)$$

where $\text{sign}(\cdot)$ is a function that returns 1 if its argument is positive and -1 if its argument is negative. Consider the total effect of ϵ_{t-1}^2 as it depends on the sign of ϵ_{t-1} ,

$$\text{Shock coefficient} = \begin{cases} \alpha_1 + \alpha_1 \gamma_1^2 + 2\alpha_1 \gamma_1 & \text{when } \epsilon_t > 0 \\ \alpha_1 + \alpha_1 \gamma_1^2 - 2\alpha_1 \gamma_1 & \text{when } \epsilon_t < 0 \end{cases} \quad (7.35)$$

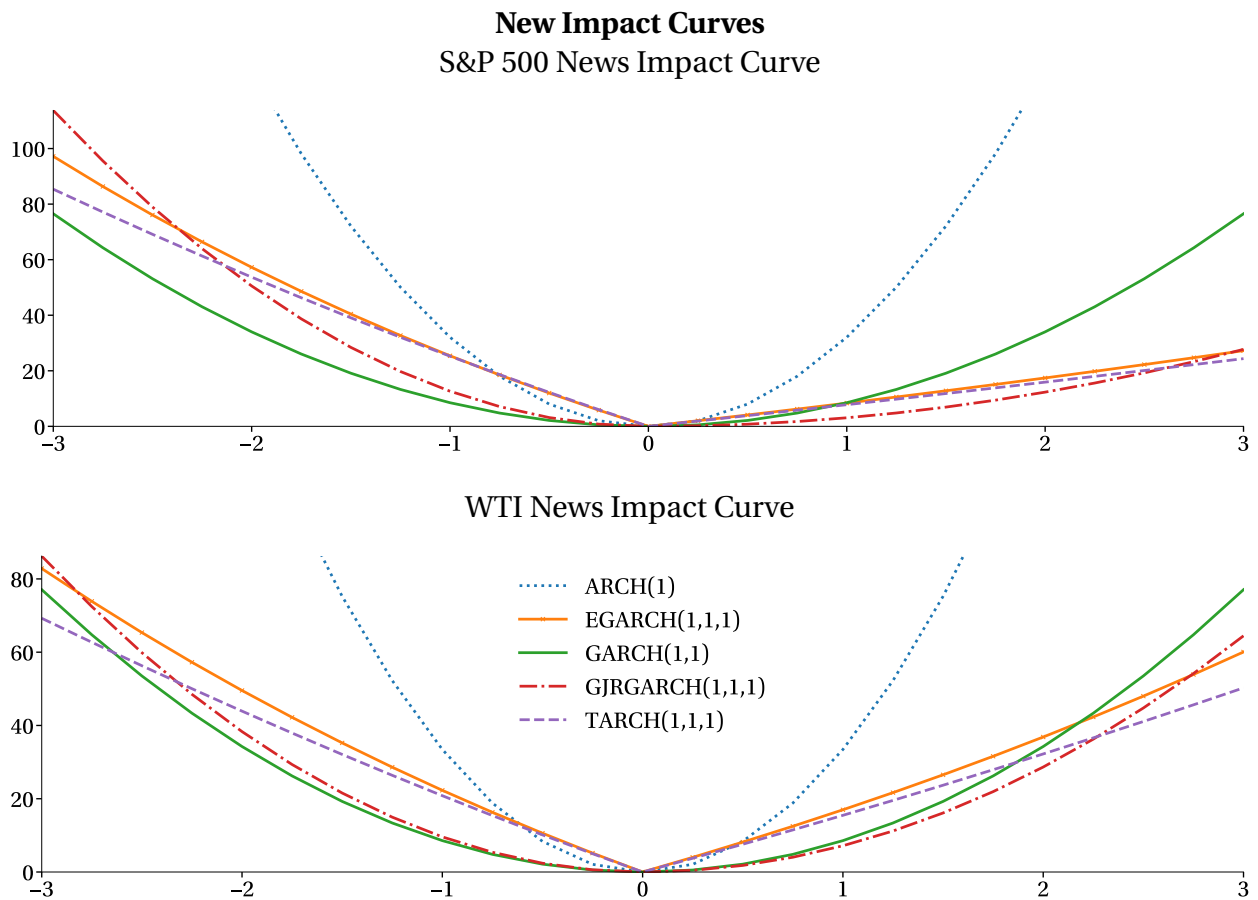


Figure 7.4: News impact curves for returns on both the S&P 500 and WTI. While the ARCH and GARCH curves are symmetric, the others show substantial asymmetries to negative news. Additionally, the fit APARCH models chose $\hat{\delta} \approx 1$, and so the NIC of the APARCH and the TARCH models appear similar.

γ is usually estimated to be less than zero which corresponds to the typical “leverage effect” in GJR-GARCH models.⁸ The relationship between a TARCH model and an APARCH model works analogously by setting $\delta = 1$. The APARCH model also nests the ARCH(P), GARCH(P,Q) and AVGARCH(P,Q) models as special cases when $\gamma_1 = 0$.

7.2.5 The News Impact Curve

With a wide range of volatility models, each with a different specification for the dynamics of conditional variances, it can be difficult to determine the precise effect of a shock to the conditional variance. *News impact curves* measure the effect of a shock in the current period on the conditional variance in the subsequent period, and so facilitate comparison between models.

⁸The explicit relationship between an APARCH and a GJR-GARCH can be derived when $\delta = 2$ by solving a system of two equations in two unknowns where eq. (7.35) is equated with the effect in a GJR-GARCH model.

Definition 7.9 (News Impact Curve (NIC)). The news impact curve of an ARCH-family model is defined as the difference between the variance with a shock e_t and the variance with no shock ($e_t = 0$). The variance in all previous periods is set to the unconditional expectation of the variance, $\bar{\sigma}^2$,

$$n(e_t) = \sigma_{t+1}^2(e_t | \sigma_t^2 = \bar{\sigma}^2) \quad (7.36)$$

$$NIC(e_t) = n(e_t) - n(0). \quad (7.37)$$

Setting the variance in the current period to the unconditional variance has two consequences. First, it ensures that the NIC does not depend on the level of variance. Second, this choice for the lagged variance improves the comparison of linear and non-linear specifications (e.g., EGARCH).

News impact curves for ARCH and GARCH models only depend on the terms which include ϵ_t^2 .

GARCH(1,1)

$$n(e_t) = \omega + \alpha_1 \bar{\sigma}^2 e_t^2 + \beta_1 \bar{\sigma}^2 \quad (7.38)$$

$$NIC(e_t) = \alpha_1 \bar{\sigma}^2 e_t^2 \quad (7.39)$$

News impact curve are more complicated when models is not linear in ϵ_t^2 . For example, consider the NIC for a TARCH(1,1,1),

$$\sigma_t = \omega + \alpha_1 |\epsilon_t| + \gamma_1 |\epsilon_t| I_{[\epsilon_t < 0]} + \beta_1 \sigma_{t-1}. \quad (7.40)$$

$$n(e_t) = \omega^2 + 2\omega(\alpha_1 + \gamma_1 I_{[\epsilon_t < 0]})|\epsilon_t| + 2\beta(\alpha_1 + \gamma_1 I_{[\epsilon_t < 0]})|\epsilon_t|\bar{\sigma} + \beta_1^2 \bar{\sigma}^2 + 2\omega\beta_1 \bar{\sigma} + (\alpha_1 + \gamma_1 I_{[\epsilon_t < 0]})^2 \epsilon_t^2 \quad (7.41)$$

$$NIC(e_t) = (\alpha_1 + \gamma_1 I_{[\epsilon_t < 0]})^2 \epsilon_t^2 + (2\omega + 2\beta_1 \bar{\sigma})(\alpha_1 + \gamma_1 I_{[\epsilon_t < 0]})|\epsilon_t| \quad (7.42)$$

While deriving explicit expressions for NICs can be tedious, practical implementation only requires computing the conditional variance for a shock of 0 ($n(0)$) and a set of shocks between -3 and 3 ($n(z)$ for $z \in (-3, 3)$). The difference between the conditional variance with a shock and the conditional variance without a shock is the NIC.

7.2.5.1 The S&P 500 and WTI

Figure 7.4 contains plots of the news impact curves for both the S&P 500 and WTI. When the models include asymmetries, the news impact curves are asymmetric and show a much larger response to negative shocks than to positive shocks, although the asymmetry is stronger in the volatility of the returns of the S&P 500 than it is in the volatility of WTI's returns.

7.3 Estimation and Inference

Consider a simple GARCH(1,1) specification,

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta \sigma_{t-1}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned} \tag{7.43}$$

Since the errors are assumed to be conditionally i.i.d. normal⁹, maximum likelihood is a natural choice to estimate the unknown parameters, θ which contain both the mean and variance parameters. The normal likelihood for T independent variables is

$$f(\mathbf{r}; \theta) = \prod_{t=1}^T (2\pi\sigma_t^2)^{-\frac{1}{2}} \exp\left(-\frac{(r_t - \mu_t)^2}{2\sigma_t^2}\right) \tag{7.44}$$

and the normal log-likelihood function is

$$l(\theta; \mathbf{r}) = \sum_{t=1}^T -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma_t^2) - \frac{(r_t - \mu_t)^2}{2\sigma_t^2}. \tag{7.45}$$

If the mean is set to 0, the log-likelihood simplifies to

$$l(\theta; \mathbf{r}) = \sum_{t=1}^T -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma_t^2) - \frac{r_t^2}{2\sigma_t^2}, \tag{7.46}$$

and is maximized by solving the first order conditions.

$$\frac{\partial l(\theta; \mathbf{r})}{\partial \sigma_t^2} = \sum_{t=1}^T -\frac{1}{2\sigma_t^2} + \frac{r_t^2}{2\sigma_t^4} = 0, \tag{7.47}$$

which can be rewritten to provide some insight into the estimation of ARCH models,

$$\frac{\partial l(\theta; \mathbf{r})}{\partial \sigma_t^2} = \frac{1}{2} \sum_{t=1}^T \frac{1}{\sigma_t^2} \left(\frac{r_t^2}{\sigma_t^2} - 1 \right). \tag{7.48}$$

This expression clarifies that the parameters of the volatility are chosen to make $\left(\frac{r_t^2}{\sigma_t^2} - 1 \right)$ as

⁹The use of conditional is to denote the dependence on σ_t^2 , which is in \mathcal{F}_{t-1} .

close to zero as possible.¹⁰ These first order conditions are not complete since ω , α_1 and β_1 , not σ_t^2 , are the parameters of a GARCH(1,1) model and

$$\frac{\partial l(\boldsymbol{\theta}; \mathbf{r})}{\partial \theta_i} = \frac{\partial l(\boldsymbol{\theta}; \mathbf{r})}{\partial \sigma_t^2} \frac{\partial \sigma_t^2}{\partial \theta_i} \quad (7.49)$$

The derivatives follow a recursive form not previously encountered,

$$\begin{aligned} \frac{\partial \sigma_t^2}{\partial \omega} &= 1 + \beta_1 \frac{\partial \sigma_{t-1}^2}{\partial \omega} \\ \frac{\partial \sigma_t^2}{\partial \alpha_1} &= \epsilon_{t-1}^2 + \beta_1 \frac{\partial \sigma_{t-1}^2}{\partial \alpha_1} \\ \frac{\partial \sigma_t^2}{\partial \beta_1} &= \sigma_{t-1}^2 + \beta_1 \frac{\partial \sigma_{t-1}^2}{\partial \beta_1}, \end{aligned} \quad (7.50)$$

although the recursion in the first order condition for ω can be removed noting that

$$\frac{\partial \sigma_t^2}{\partial \omega} = 1 + \beta_1 \frac{\partial \sigma_{t-1}^2}{\partial \omega} \approx \frac{1}{1 - \beta_1}. \quad (7.51)$$

Eqs. (7.49) – (7.51) provide the necessary formulas to implement the scores of the log-likelihood although they are not needed to estimate a GARCH model.¹¹

The use of the normal likelihood has one strong justification; estimates produced by maximizing the log-likelihood of a normal are *strongly consistent*. Strong consistency is a property of an estimator that ensures parameter estimates converge to the true parameters *even if the assumed conditional distribution is misspecified*. For example, in a standard GARCH(1,1), the parameter estimates would still converge to their true value if estimated with the normal likelihood as long as the volatility model was correctly specified. The intuition behind this result comes from the *generalized error*

$$\left(\frac{\epsilon_t^2}{\sigma_t^2} - 1 \right). \quad (7.52)$$

Whenever $\sigma_t^2 = E_{t-1}[\epsilon_t^2]$,

¹⁰If $E_{t-1} \left[\frac{r_t^2}{\sigma_t^2} - 1 \right] = 0$, and so the volatility is correctly specified, then the scores of the log-likelihood have expectation zero since

$$\begin{aligned} E \left[\frac{1}{\sigma_t^2} \left(\frac{r_t^2}{\sigma_t^2} - 1 \right) \right] &= E \left[E_{t-1} \left[\frac{1}{\sigma_t^2} \left(\frac{r_t^2}{\sigma_t^2} - 1 \right) \right] \right] \\ &= E \left[\frac{1}{\sigma_t^2} \left(E_{t-1} \left[\frac{r_t^2}{\sigma_t^2} - 1 \right] \right) \right] \\ &= E \left[\frac{1}{\sigma_t^2} (0) \right] \\ &= 0. \end{aligned}$$

¹¹MATLAB and many other econometric packages are capable of estimating the derivatives using a numerical approximation that only requires the log-likelihood. Numerical derivatives use the definition of a derivative, $f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$ to approximate the derivative using $f'(x) \approx \frac{f(x+h) - f(x)}{h}$ for some small h .

$$E \left[\left(\frac{\epsilon_t^2}{\sigma_t^2} - 1 \right) \right] = E \left[\left(\frac{E_{t-1}[\epsilon_t^2]}{\sigma_t^2} - 1 \right) \right] = E \left[\left(\frac{\sigma_t^2}{\sigma_t^2} - 1 \right) \right] = 0. \quad (7.53)$$

Thus, as long as the GARCH model nests the true DGP, the parameters are chosen to make the conditional expectation of the generalized error 0; these parameters correspond to those of the original DGP even if the conditional distribution is misspecified.¹² This is a unique property of the normal distribution and is not found in other common distributions.

7.3.1 Inference

Under some regularity conditions, parameters estimated by maximum likelihood are asymptotically normally distributed,

$$\sqrt{T}(\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}_0) \xrightarrow{d} N(0, \mathcal{I}^{-1}) \quad (7.54)$$

where

$$\mathcal{I} = -E \left[\frac{\partial^2 l(\boldsymbol{\theta}_0; r_t)}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}'} \right] \quad (7.55)$$

is the negative of the expected Hessian. The Hessian measures how much curvature there is in the log-likelihood at the optimum just like the second-derivative measures the rate-of-change in the rate-of-change of the function in a standard calculus problem. The sample analog estimator that averages the time-series of Hessian matrices computed at $\hat{\boldsymbol{\theta}}$ is used to estimate \mathcal{I} ,

$$\hat{\mathcal{I}} = -T^{-1} \sum_{t=1}^T \frac{\partial^2 l(\hat{\boldsymbol{\theta}}; r_t)}{\partial \boldsymbol{\theta} \partial \boldsymbol{\theta}'}. \quad (7.56)$$

Chapter 2 shows that the Information Matrix Equality (IME) generally holds for MLE problems, so that

$$\mathcal{I} = \mathcal{J} \quad (7.57)$$

where

$$\mathcal{J} = E \left[\frac{\partial l(r_t; \boldsymbol{\theta}_0)}{\partial \boldsymbol{\theta}} \frac{\partial l(r_t; \boldsymbol{\theta}_0)}{\partial \boldsymbol{\theta}'} \right] \quad (7.58)$$

is the covariance of the scores. The scores behave like errors in ML estimators and so large score variance indicate the parameters are difficult to estimate accurately. The estimator of \mathcal{J} is the sample analog averaging the outer-product of the scores evaluated at the estimated parameters,

$$\hat{\mathcal{J}} = T^{-1} \sum_{t=1}^T \frac{\partial l(\hat{\boldsymbol{\theta}}; r_t)}{\partial \boldsymbol{\theta}} \frac{\partial l(\hat{\boldsymbol{\theta}}; r_t)}{\partial \boldsymbol{\theta}'}. \quad (7.59)$$

The IME generally applies when the parameter estimates are *maximum likelihood estimates*, which requires that both the likelihood used in estimation is correct and that the specification

¹²An assumption that a GARCH specification nests the DGP is extremely strong and likely wrong in most cases. However, the strong consistency property of the normal likelihood in volatility models justifies estimation of models where the standardized residuals are leptokurtotic and skewed.

for the conditional variance is general enough to nest the true process. When one specification is used for estimation (e.g., normal) but the data follow a different conditional distribution, these estimators are known as Quasi-Maximum Likelihood Estimators (QMLE), and the IME generally fails to hold. Under some regularity conditions, the estimated parameters are still asymptotically normal but with a different covariance,

$$\sqrt{T}(\hat{\boldsymbol{\theta}} - \boldsymbol{\theta}_0) \xrightarrow{d} N(0, \mathcal{I}^{-1} \mathcal{J} \mathcal{I}^{-1}) \quad (7.60)$$

If the IME is valid, $\mathcal{I} = \mathcal{J}$ and so this covariance simplifies to the usual MLE variance estimator.

In most applications of ARCH models, the conditional distribution of shocks is decidedly not normal, and standardized residuals have excess kurtosis and are skewed. Bollerslev and Wooldridge (1992) were the first to show that the IME does not generally hold for GARCH models when the distribution is misspecified and the “sandwich” form

$$\hat{\mathcal{I}}^{-1} \hat{\mathcal{J}} \hat{\mathcal{I}}^{-1} \quad (7.61)$$

of the covariance estimator is often referred to as the *Bollerslev-Wooldridge* covariance matrix or alternatively a robust covariance matrix. Standard Wald tests can be used to test hypotheses of interest, such as whether an asymmetric term is statistically significant, although likelihood ratio tests are not reliable since they do not have the usual χ_m^2 distribution.

7.3.1.1 The S&P 500 and WTI

A TARCH(1,1,1) models were estimated on both the S&P 500 and WTI returns to illustrate the differences between the MLE and the Bollerslev-Wooldridge (QMLE) covariance estimators. Table 7.3 contains the estimated parameters and t-stats using both the MLE covariance matrix and the Bollerslev-Wooldridge covariance matrix. The robust t-stats are substantially smaller than conventional ones, although conclusions about statistical significance are not affected except for ω in the WTI model. These changes are due to the heavy-tail in the standardized residuals, $\hat{\varepsilon}_t = r_t - \hat{\mu}_t / \hat{\sigma}_t$, in these series.

7.3.1.2 Independence of the mean and variance parameters

Inference on the parameters of the ARCH model is still valid when using normal MLE or QMLE when the model for the mean is general enough to nest the true form. This property is important in practice since mean and variance parameters can be estimated separately without correcting the covariance matrix of the estimated parameters.¹³ This surprising feature of QMLE estimators employing a normal log-likelihood comes from the cross-partial derivative of the log-likelihood with respect to the mean and variance parameters,

$$l(\boldsymbol{\theta}; r_t) = -\frac{1}{2} \log(2\pi) - \frac{1}{2} \log(\sigma_t^2) - \frac{(r_t - \mu_t)^2}{2\sigma_t^2}. \quad (7.62)$$

¹³The estimated covariance for the mean should use a White covariance estimator. If the mean parameters are of particular interest, it may be more efficient to jointly estimate the parameters of the mean and volatility equations as a form of GLS (see Chapter 3).

WTI				
	ω	α_1	γ_1	β_1
Coefficient	0.031	0.030	0.055	0.942
Std. T-stat	3.62	4.03	7.67	102.94
Robust T-stat	1.85	2.31	4.45	49.66

S&P 500				
	ω	α_1	γ_1	β_1
Coefficient	0.026	0.000	0.172	0.909
Std. T-stat	9.63	0.000	14.79	124.92
Robust T-stat	6.28	0.000	10.55	93.26

Table 7.3: Estimates from a TARCH(1,1,1) for the S&P 500 and WTI using alternative parameter covariance estimators.

The first order condition is,

$$\frac{\partial l(\boldsymbol{\theta}; \mathbf{r})}{\partial \mu_t} \frac{\partial \mu_t}{\partial \phi} = - \sum_{t=1}^T \frac{(r_t - \mu_t)}{\sigma_t^2} \frac{\partial \mu_t}{\partial \phi} \quad (7.63)$$

and the second order condition is

$$\frac{\partial^2 l(\boldsymbol{\theta}; \mathbf{r})}{\partial \mu_t \partial \sigma_t^2} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} = \sum_{t=1}^T \frac{(r_t - \mu_t)}{\sigma_t^4} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \quad (7.64)$$

where ϕ is a parameter of the conditional mean and ψ is a parameter of the conditional variance. For example, in a simple ARCH(1) model with a constant mean,

$$\begin{aligned} r_t &= \mu + \epsilon_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1), \end{aligned} \quad (7.65)$$

$\phi = \mu$ and ψ can be either ω or α_1 . Taking expectations of the cross-partial,

$$\begin{aligned} E \left[\frac{\partial^2 l(\boldsymbol{\theta}; \mathbf{r})}{\partial \mu_t \partial \sigma_t^2} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \right] &= E \left[\sum_{t=1}^T \frac{r_t - \mu_t}{\sigma_t^4} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \right] \\ &= E \left[E_{t-1} \left[\sum_{t=1}^T \frac{r_t - \mu_t}{\sigma_t^4} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \right] \right] \end{aligned} \quad (7.66)$$

$$\begin{aligned}
&= E \left[\sum_{t=1}^T \frac{E_{t-1} [r_t - \mu_t]}{\sigma_t^4} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \right] \\
&= E \left[\sum_{t=1}^T \frac{0}{\sigma_t^4} \frac{\partial \mu_t}{\partial \phi} \frac{\partial \sigma_t^2}{\partial \psi} \right] \\
&= 0
\end{aligned}$$

it can be seen that the expectation of the cross derivative is 0. The intuition behind the result follows from noticing that when the mean model is correct for the conditional expectation of r_t , the term $r_t - \mu_t$ has conditional expectation 0 and knowledge of the variance is not needed. This argument is a similar one used to establish the validity of the OLS estimator when the errors are heteroskedastic.

7.4 GARCH-in-Mean

The GARCH-in-mean model (GiM) makes a significant change to the role of time-varying volatility by explicitly relating the level of volatility to the expected return (Engle, Lilien, and Robins, 1987). A simple GiM model can be specified as

$$\begin{aligned}
r_t &= \mu + \delta \sigma_t^2 + \epsilon_t \\
\sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2 \\
\epsilon_t &= \sigma_t e_t \\
e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1)
\end{aligned} \tag{7.67}$$

although virtually any ARCH-family model could be used to model the conditional variance. The obvious difference between the GiM and a standard GARCH(1,1) is that the variance appears in the mean of the return. Note that the shock driving the changes in variance is not the mean return but still ϵ_{t-1}^2 , and so the ARCH portion of a GiM is unaffected. Other forms of the GiM model have been employed where the conditional standard deviation or the log of the conditional variance are used in the mean equation¹⁴,

$$r_t = \mu + \delta \sigma_t + \epsilon_t \tag{7.68}$$

or

$$r_t = \mu + \delta \ln(\sigma_t^2) + \epsilon_t \tag{7.69}$$

Because the variance appears in the mean equation for r_t , the mean and variance parameters cannot be separately estimated. Despite the apparent feedback, processes that follow a GiM are stationary as long as the variance process is stationary. The conditional variance (σ_t^2) in the conditional mean does not feedback into the conditional variance process and so behaves like an exogenous regressor.

¹⁴The model for the conditional mean can be extended to include ARMA terms or any other predetermined regressor.

S&P 500 Garch-in-Mean Estimates							
	μ	δ	ω	α	γ	β	Log Lik.
σ^2	0.004 (0.753)	0.022 (0.074)	0.022 (0.000)	0.000 (0.999)	0.183 (0.000)	0.888 (0.000)	−6773.7
σ	−0.034 (0.304)	0.070 (0.087)	0.022 (0.000)	0.000 (0.999)	0.182 (0.000)	0.887 (0.000)	−6773.4
$\ln \sigma^2$	0.038 (0.027)	0.030 (0.126)	0.022 (0.000)	0.000 (0.999)	0.183 (0.000)	0.888 (0.000)	−6773.8

Table 7.4: GARCH-in-mean estimates for the S&P 500 series. δ measures the strength of the GARCH-in-mean, and so is the most interesting parameter. The volatility process was a standard GARCH(1,1). P-values are in parentheses.

7.4.1 The S&P 500

Standard asset pricing theory dictates that there is a risk-return trade-off. GARCH-in-mean models provide a natural method to test whether this is the case. Using the S&P 500 data, three GiM models were estimated (one for each transformation of the variance in the mean equation), and the results are presented in table 7.4. Based on these estimates, there does appear to be a trade-off between mean and variance and higher variances produce higher expected means, although the magnitude is economically small and the coefficients are only significant at the 10% level.

7.5 Alternative Distributional Assumptions

Despite the strengths of the assumption that the errors are conditionally normal – estimation is simple, and parameters are *strongly consistent* for the true parameters – GARCH models can be specified and estimated with alternative distributional assumptions. The motivation for using something other than the normal distribution is two-fold. First, a better approximation to the conditional distribution of the standardized returns may improve the precision of the volatility process parameter estimates and, in the case of MLE, the estimates will be fully efficient. Second, GARCH models are often used in applications where the choice of the assumed density is plays a larger role such as in Value-at-Risk estimation or option pricing.

Three distributions stand among the dozens that have been used to estimate the parameters of GARCH processes. The first is a standardized Student's t (to have a unit variance for any value ν , see Bollerslev (1987)) with ν degrees of freedom,

Standardized Student's t

$$f(r_t; \mu, \sigma_t^2, \nu) = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\Gamma\left(\frac{\nu}{2}\right)} \frac{1}{\sqrt{\pi(\nu-2)}} \frac{1}{\sigma_t} \frac{1}{\left(1 + \frac{(r_t - \mu)^2}{\sigma_t^2(\nu-2)}\right)^{\frac{\nu+1}{2}}} \quad (7.70)$$

where $\Gamma(\cdot)$ is the gamma function.¹⁵ This distribution is always fat-tailed and produces a better

¹⁵The standardized Student's t differs from the usual Student's t so that it is necessary to scale data by $\sqrt{\frac{\nu}{\nu-2}}$ if using functions (such as the CDF) for the regular Student's t distribution.

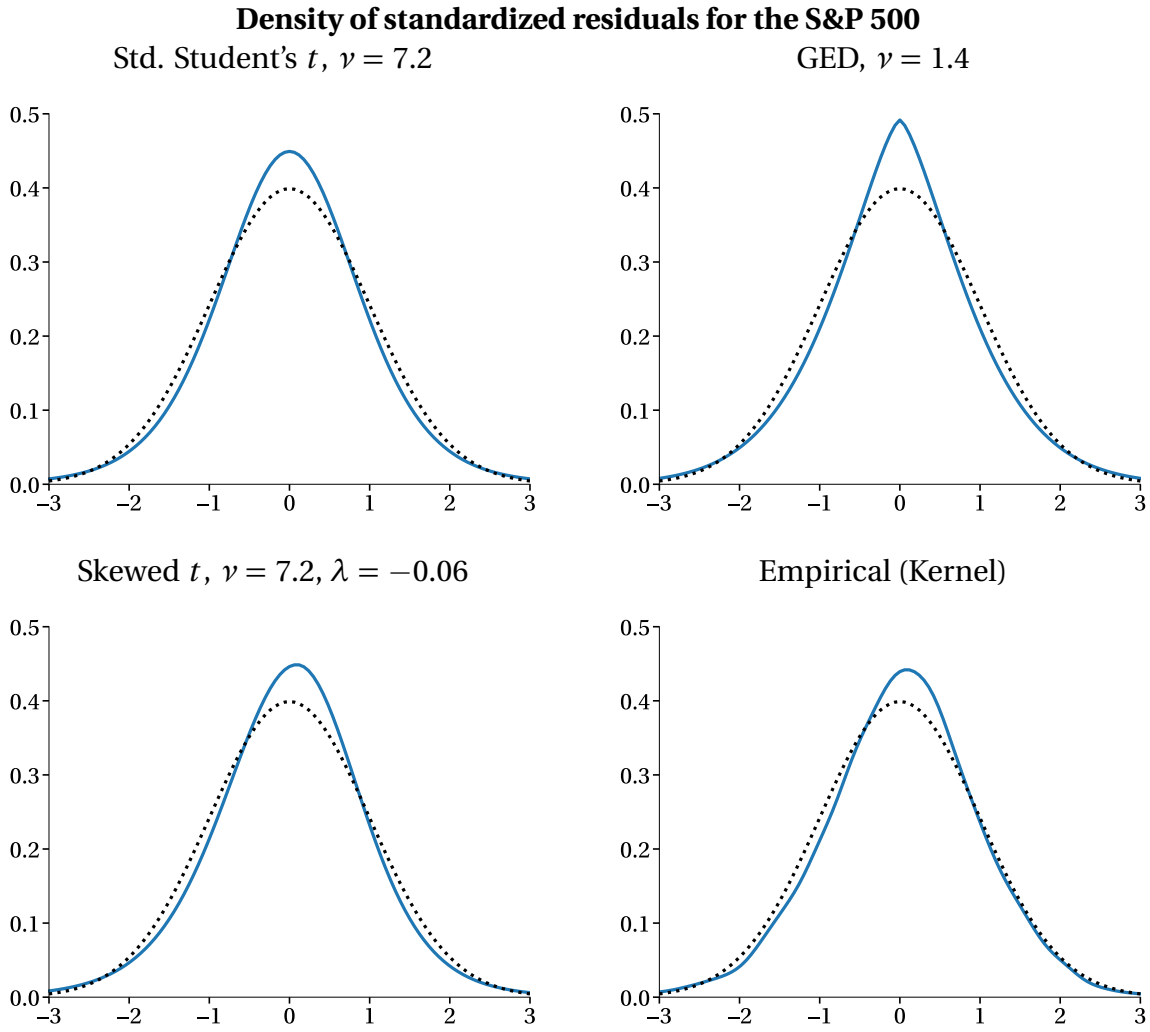


Figure 7.5: The four panels of this figure contain the estimated density for the S&P 500 and the density implied by the distributions: Student's t , GED, Hansen's Skew t and a kernel density plot of the standardized residuals, $\hat{\epsilon}_t = \epsilon_t / \hat{\sigma}_t$, along with the PDF of a normal (dotted line) for comparison. The shape parameters, ν and λ , were jointly estimated with the variance parameters in the Student's t , GED, and skewed t .

fit than the normal for most asset return series. This distribution is only well defined if $\nu > 2$ since the variance of a Student's t with $\nu \leq 2$ is infinite. The second is the generalized error distribution (GED, see Nelson (1991)),

Generalized Error Distribution

$$f(r_t; \mu, \sigma_t^2, \nu) = \frac{\nu \exp\left(-\frac{1}{2} \left| \frac{r_t - \mu}{\sigma_t \lambda} \right|^\nu\right)}{\sigma_t \lambda 2^{\frac{\nu+1}{\nu}} \Gamma(\frac{1}{\nu})} \quad (7.71)$$

$$\lambda = \sqrt{\frac{2^{-\frac{2}{\nu}} \Gamma(\frac{1}{\nu})}{\Gamma(\frac{3}{\nu})}} \quad (7.72)$$

which nests the normal when $\nu = 2$. The GED is fat-tailed when $\nu < 2$ and thin-tailed when $\nu > 2$. It is necessary that $\nu \geq 1$ to use the GED in volatility model estimation to ensure that variance is finite. The third useful distribution, introduced in Hansen (1994), extends the standardized Student's t to allow for skewness of returns

Hansen's skewed t

$$f(\epsilon_t; \mu, \sigma_t, \nu, \lambda) = \begin{cases} bc \left(1 + \frac{1}{\nu-2} \left(\frac{b \left(\frac{r_t - \mu}{\sigma_t} \right) + a}{(1-\lambda)} \right)^2 \right)^{-(\nu+1)/2}, & \frac{r_t - \mu}{\sigma_t} < -a/b \\ bc \left(1 + \frac{1}{\nu-2} \left(\frac{b \left(\frac{r_t - \mu}{\sigma_t} \right) + a}{(1+\lambda)} \right)^2 \right)^{-(\nu+1)/2}, & \frac{r_t - \mu}{\sigma_t} \geq -a/b \end{cases} \quad (7.73)$$

where

$$a = 4\lambda c \left(\frac{\nu-2}{\nu-1} \right),$$

$$b = \sqrt{1 + 3\lambda^2 - a^2},$$

and

$$c = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\sqrt{\pi(\nu-2)}\Gamma\left(\frac{\nu}{2}\right)}.$$

The two shape parameters, ν and λ , control the kurtosis and the skewness, respectively.

These distributions may be better approximations to the actual distribution of the standardized residuals since they allow for kurtosis greater than that of the normal, an important empirical fact, and, in the case of the skewed t , skewness in the standardized returns. Chapter 8 applies these distributions in the context of Value-at-Risk and density forecasting.


7.5.1 Alternative Distribution in Practice

To explore the role of alternative distributional assumptions in the estimation of GARCH models, a TARCH(1,1,1) was fit to the S&P 500 returns using the conditional normal, the Student's t , the GED and Hansen's skewed t . Figure 7.5 contains the empirical density (constructed with a kernel) and the fit density of the three distributions. The shape parameters, ν and λ , were jointly estimated with the conditional variance parameters. Figure 7.6 plots of the estimated conditional variance for both the S&P 500 and WTI using all four distributional assumptions. The most important aspect of this figure is that the fit variances are indistinguishable. This is a common finding: estimating models using alternative distributional assumptions produce little difference in the estimated parameters or the fitted conditional variances from the volatility model.¹⁶

¹⁶While the volatilities are similar, the models do not fit the data equally well. The alternative distributions often provide a better fit as measured by the log-likelihood and provide a more accurate description of the probability in the tails of the distribution.

Conditional Variance and Distributional Assumptions

S&P 500 Annualized Volatility (TARCH(1,1,1))



WTI Annualized Volatility (TARCH(1,1,1))




Figure 7.6: The choice of the distribution for the standardized innovation makes little difference to the fit variances or the estimated parameters in most models. The alternative distributions are more useful in application to Value-at-Risk and Density forecasting where the choice of density plays a more significant role.

7.6 Model Building

Since ARCH and GARCH models are similar to AR and ARMA models, the Box-Jenkins methodology is a natural way to approach the problem. The first step is to analyze the sample ACF and PACF of the *squared* returns, or if the model for the conditional mean is non-trivial, the sample ACF and PACF of the estimated residuals, $\hat{\epsilon}_t$, should be examined for heteroskedasticity. Figures 7.7 and 7.8 contains the ACF and PACF for the squared returns of the S&P 500 and WTI respectively. The models used in selecting the final model are reproduced in tables 7.5 and 7.6 respectively. Both selections began with a simple GARCH(1,1). The next step was to check if more lags were needed for either the squared innovation or the lagged variance by fitting a GARCH(2,1) and a GARCH(1,2) to each series. Neither of these meaningfully improved the fit, and a GARCH(1,1) was assumed to be sufficient to capture the symmetric dynamics.

The next step in model building is to examine whether the data exhibit any evidence of asymmetries using a GJR-GARCH(1,1,1). The asymmetry term was significant and so other forms of

	α_1	α_2	γ_1	γ_2	β_1	β_2	Log Lik.
GARCH(1,1)	0.102 (0.000)				0.885 (0.000)		−6887.6
GARCH(1,2)	0.102 (0.000)				0.885 (0.000)	0.000 (0.999)	−6887.6
GARCH(2,1)	0.067 (0.003)	0.053 (0.066)			0.864 (0.000)		−6883.5
GJR-GARCH(1,1,1)	0.000 (0.999)		0.185 (0.000)		0.891 (0.000)		−6775.1
GJR-GARCH(1,2,1)	0.000 (0.999)		0.158 (0.000)	0.033 (0.460)	0.887 (0.000)		−6774.5
TARCH(1,1,1)*	0.000 (0.999)		0.172 (0.000)		0.909 (0.000)		−6751.9
TARCH(1,2,1)	0.000 (0.999)		0.165 (0.000)	0.009 (0.786)	0.908 (0.000)		−6751.8
TARCH(2,1,1)	0.000 (0.999)	0.003 (0.936)	0.171 (0.000)		0.907 (0.000)		−6751.9
EGARCH(1,0,1)	0.211 (0.000)				0.979 (0.000)		−6908.4
EGARCH(1,1,1)	0.136 (0.000)		−0.153 (0.000)		0.975 (0.000)		−6766.7
EGARCH(1,2,1)	0.129 (0.000)		−0.213 (0.000)	0.067 (0.045)	0.977 (0.000)		−6761.7
EGARCH(2,1,1)	0.020 (0.651)	0.131 (0.006)	−0.162 (0.000)		0.970 (0.000)		−6757.6

Table 7.5: The models estimated in selecting a final model for the conditional variance of the S&P 500 Index. * indicates the selected model.

the GJR model were explored. All were found to provide little improvement in the fit. Once a GJR-GARCH(1,1,1) model was decided upon, a TARCH(1,1,1) was fit to examine whether evolution in variances or standard deviations was more appropriate for the data. Both series preferred the TARCH to the GJR-GARCH (compare the log-likelihoods), and the TARCH(1,1,1) was selected. In comparing alternative specifications, an EGARCH was fit and found to provide a good description of the data. In both cases, the EGARCH was expanded to include more lags of the shocks or lagged log volatility. The EGARCH did not improve over the TARCH for the S&P 500, and so the TARCH(1,1,1) was selected. The EGARCH did fit the WTI data better, and so the preferred model is an EGARCH(1,1,1), although a case could be made for the EGARCH(2,1,1) which provided a better fit. Overfitting is always a concern, and the opposite signs on α_1 and α_2 in the EGARCH(2,1,1) are suspicious.

7.6.0.1 Testing for (G)ARCH

Although conditional heteroskedasticity can often be identified by graphical inspection, a formal test of conditional homoskedasticity is also useful. The standard method to test for ARCH is to use the ARCH-LM test which is implemented as a regression of *squared* residuals on lagged squared residuals. The test directly exploits the AR representation of an ARCH process (Engle, 1982) and is computed as T times the R^2 ($LM = T \times R^2$) from the regression

$$\hat{\epsilon}_t^2 = \phi_0 + \phi_1 \hat{\epsilon}_{t-1}^2 + \dots + \phi_P \hat{\epsilon}_{t-P}^2 + \eta_t. \quad (7.74)$$

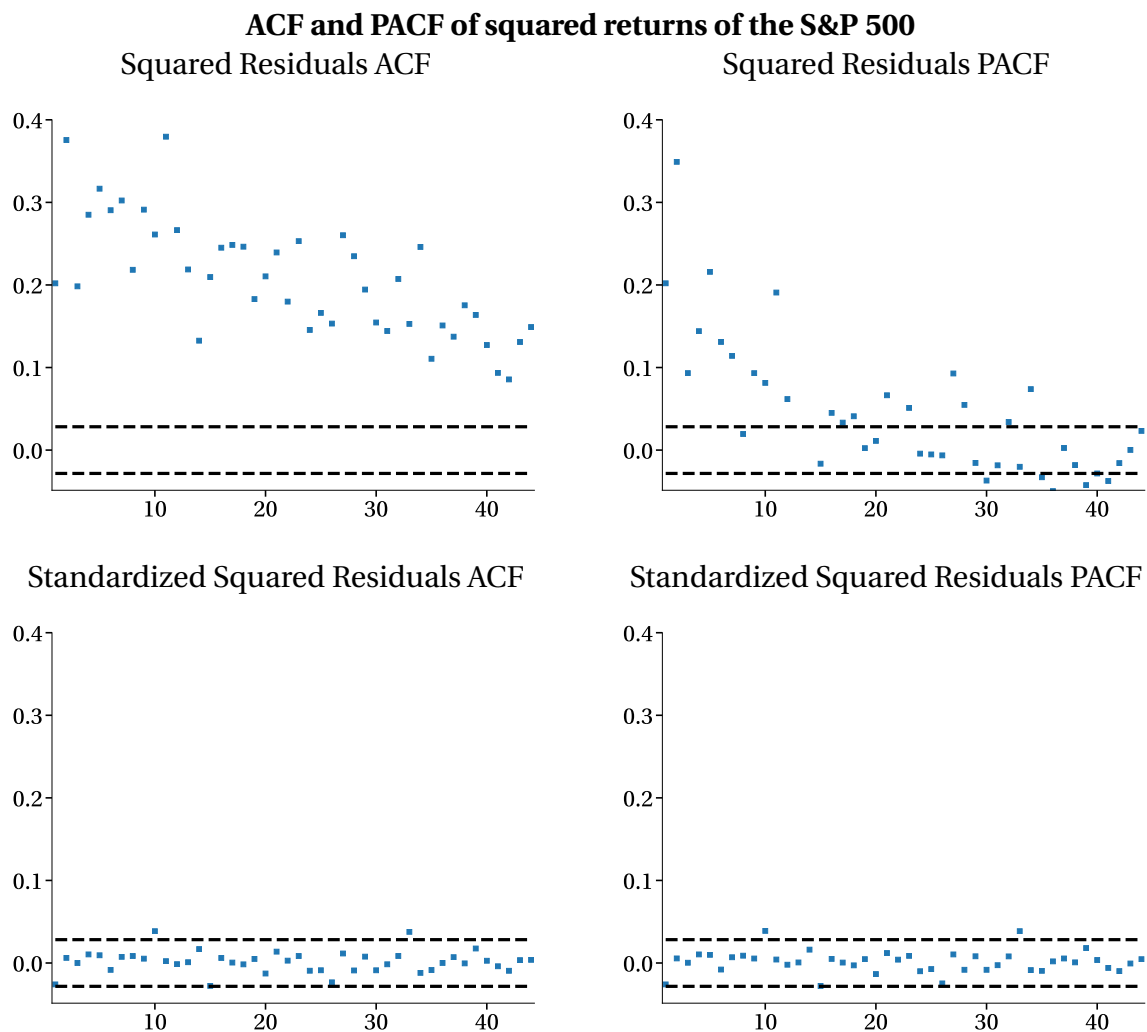


Figure 7.7: ACF and PACF of the squared returns for the S&P 500. The bottom two panels plot the ACF and PACF of the standardized squared residuals, $\hat{e}_t^2 = \hat{e}_t^2 / \hat{\sigma}_t^2$. The top panels indicate persistence through both the ACF and PACF. These plots suggest that a GARCH model is needed. The ACF and PACF of the standardized residuals are consistent with those of a white noise process.

The test statistic is asymptotically distributed χ_p^2 where \hat{e}_t are residuals constructed from the returns by subtracting the conditional mean. The null hypothesis is $H_0 : \phi_1 = \dots = \phi_p = 0$ which corresponds to no persistence in the conditional variance.

7.7 Forecasting Volatility

Forecasting conditional variances with ARCH-family models ranges from simple for ARCH and GARCH processes to difficult for non-linear specifications. Consider the simple ARCH(1) pro-

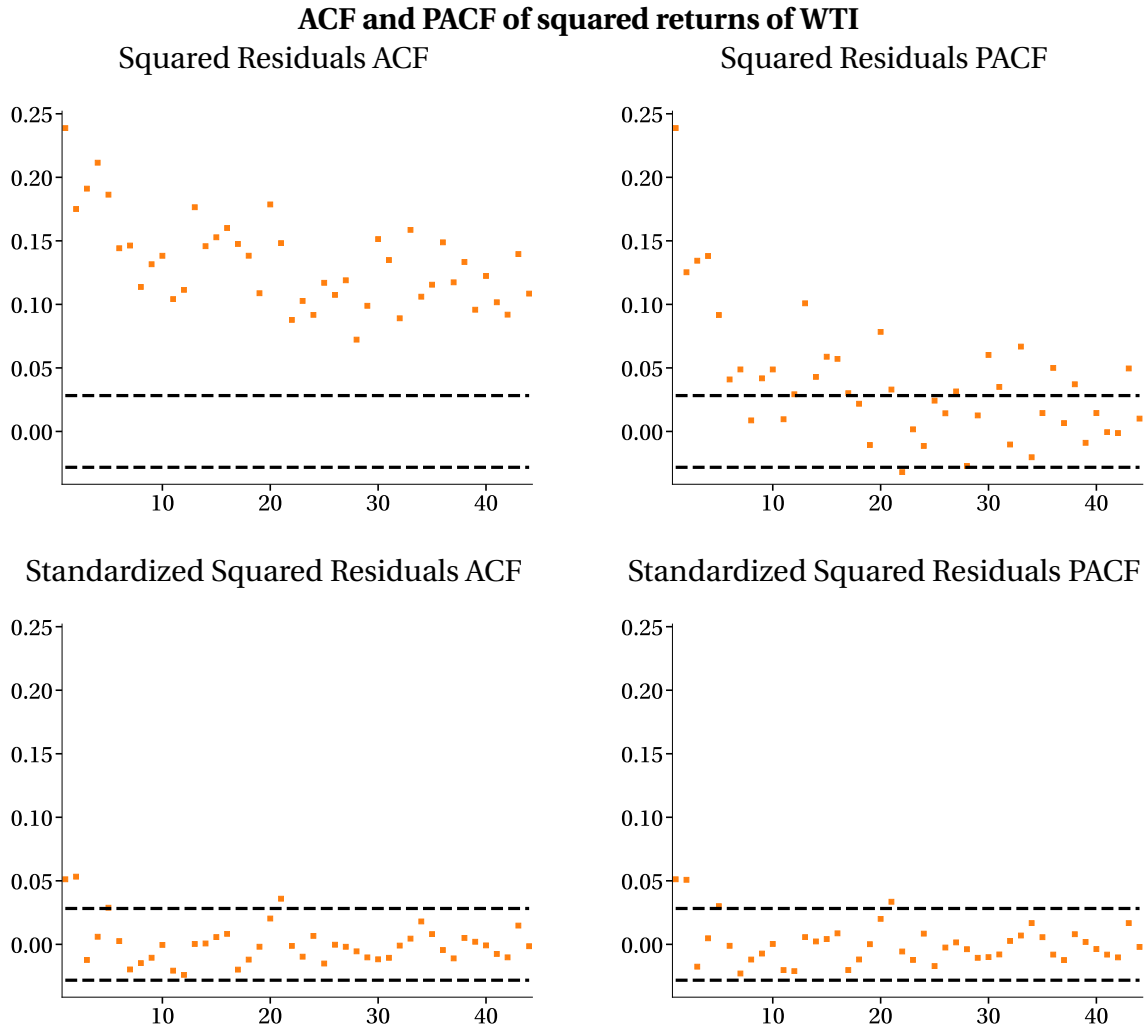


Figure 7.8: ACF and PACF of the squared returns for WTI. The bottom two panels plot the ACF and PACF of the standardized squared residuals, $\hat{\epsilon}_t^2 = \hat{\epsilon}_t^2 / \hat{\sigma}_t^2$. The top panels indicate persistence through both the ACF and PACF. These plots suggest that a GARCH model is needed. The ACF and PACF of the standardized residuals are consistent with those of a white noise process. When compared to the S&P 500 ACF and PACF, the ACF and PACF of the WTI returns indicate less persistence in volatility.

cess,

$$\epsilon_t = \sigma_t e_t \quad (7.75)$$

$$e_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$$

$$\sigma_t^2 = \omega + \alpha_1 \epsilon_{t-1}^2$$

Iterating forward, $\sigma_{t+1}^2 = \omega + \alpha_1 \epsilon_t^2$, and taking conditional expectations, $E_t[\sigma_{t+1}^2] = E_t[\omega + \alpha_1 \epsilon_t^2] = \omega + \alpha_1 \epsilon_t^2$ since all of these quantities are known at time t . This is a property common to

	α_1	α_2	γ_1	γ_2	β_1	β_2	Log Lik.
GARCH(1,1)	0.059 (0.000)				0.934 (0.000)		-11030.1
GARCH(1,2)	0.075 (0.000)				0.585 (0.000)	0.331 (0.027)	-11027.4
GARCH(2,1)	0.059 (0.001)	0.000 (0.999)			0.934 (0.000)		-11030.1
GJR-GARCH(1,1,1)	0.026 (0.008)		0.049 (0.000)		0.945 (0.000)		-11011.9
GJR-GARCH(1,2,1)	0.026 (0.010)		0.049 (0.102)	0.000 (0.999)	0.945 (0.000)		-11011.9
TARCH(1,1,1)	0.030 (0.021)		0.055 (0.000)		0.942 (0.000)		-11005.6
TARCH(1,2,1)	0.030 (0.038)		0.055 (0.048)	0.000 (0.999)	0.942 (0.000)		-11005.6
TARCH(2,1,1)	0.030 (0.186)	0.000 (0.999)	0.055 (0.000)		0.942 (0.000)		-11005.6
EGARCH(1,0,1)	0.148 (0.000)				0.986 (0.000)		-11029.5
EGARCH(1,1,1) [†]	0.109 (0.000)		-0.050 (0.000)		0.990 (0.000)		-11000.6
EGARCH(1,2,1)	0.109 (0.000)		-0.056 (0.043)	0.006 (0.834)	0.990 (0.000)		-11000.5
EGARCH(2,1,1) [*]	0.195 (0.000)	-0.101 (0.019)	-0.049 (0.000)		0.992 (0.000)		-10994.4

Table 7.6: The models estimated in selecting a final model for the conditional variance of WTI. ^{*} indicates the selected model. [†] indicates a model that could be considered for model selection.

all ARCH-family models: the forecast of σ_{t+1}^2 is known at time t .¹⁷

The 2-step ahead forecast follows from an application of the law of iterated expectations,

$$\begin{aligned}
 E_t[\sigma_{t+2}^2] &= E_t[\omega + \alpha_1 \epsilon_{t+1}^2]. \\
 &= \omega + \alpha_1 E_t[\epsilon_{t+1}^2] \\
 &= \omega + \alpha_1 (\omega + \alpha_1 \epsilon_t^2) \\
 &= \omega + \alpha_1 \omega + \alpha_1^2 \epsilon_t^2
 \end{aligned} \tag{7.76}$$

The expression for an h -step ahead forecast can be constructed by repeated substitution and is given by

$$E_t[\sigma_{t+h}^2] = \sum_{i=0}^{h-1} \alpha_1^i \omega + \alpha_1^h \epsilon_t^2. \tag{7.77}$$

An ARCH(1) is an AR(1), and this formula is identical to the expression for the multi-step forecast of an AR(1).

Forecasts from GARCH(1,1) models are constructed following the same steps. The one-step-ahead forecast is

¹⁷Not only is this property common to all ARCH-family members, but it is also the defining characteristic of an ARCH model.

$$\begin{aligned} E_t[\sigma_{t+1}^2] &= E_t[\omega + \alpha_1 \epsilon_t^2 + \beta_1 \sigma_t^2] \\ &= \omega + \alpha_1 \epsilon_t^2 + \beta_1 \sigma_t^2. \end{aligned} \quad (7.78)$$

The two-step-ahead forecast is

$$\begin{aligned} E_t[\sigma_{t+2}^2] &= E_t[\omega + \alpha_1 \epsilon_{t+1}^2 + \beta_1 \sigma_{t+1}^2] \\ &= \omega + \alpha_1 E_t[\epsilon_{t+1}^2] + \beta_1 E_t[\sigma_{t+1}^2] \\ &= \omega + \alpha_1 E_t[e_{t+1}^2 \sigma_{t+1}^2] + \beta_1 E_t[\sigma_{t+1}^2] \\ &= \omega + \alpha_1 E_t[e_{t+1}^2] E_t[\sigma_{t+1}^2] + \beta_1 E_t[\sigma_{t+1}^2] \\ &= \omega + \alpha_1 \cdot 1 \cdot E_t[\sigma_{t+1}^2] + \beta_1 E_t[\sigma_{t+1}^2] \\ &= \omega + \alpha_1 E_t[\sigma_{t+1}^2] + \beta_1 E_t[\sigma_{t+1}^2] \\ &= \omega + (\alpha_1 + \beta_1) E_t[\sigma_{t+1}^2]. \end{aligned}$$

Substituting the one-step-ahead forecast, $E_t[\sigma_{t+1}^2]$, shows that the forecast only depends on time t information,

$$\begin{aligned} E_t[\sigma_{t+2}^2] &= \omega + (\alpha_1 + \beta_1)(\omega + \alpha_1 \epsilon_t^2 + \beta_1 \sigma_t^2) \\ &= \omega + (\alpha_1 + \beta_1)\omega + (\alpha_1 + \beta_1)\alpha_1 \epsilon_t^2 + (\alpha_1 + \beta_1)\beta_1 \sigma_t^2. \end{aligned} \quad (7.79)$$

Note that $E_t[\sigma_{t+3}^2] = \omega + (\alpha_1 + \beta_1)E_t[\sigma_{t+2}^2]$, and so

$$\begin{aligned} E_t[\sigma_{t+3}^2] &= \omega + (\alpha_1 + \beta_1)(\omega + (\alpha_1 + \beta_1)\omega + (\alpha_1 + \beta_1)\alpha_1 \epsilon_t^2 + (\alpha_1 + \beta_1)\beta_1 \sigma_t^2) \\ &= \omega + (\alpha_1 + \beta_1)\omega + (\alpha_1 + \beta_1)^2 \omega + (\alpha_1 + \beta_1)^2 \alpha_1 \epsilon_t^2 + (\alpha_1 + \beta_1)^2 \beta_1 \sigma_t^2. \end{aligned} \quad (7.80)$$

Repeated substitution reveals a pattern in the multi-step forecasts which is compactly expressed as

$$E_t[\sigma_{t+h}^2] = \sum_{i=0}^{h-1} (\alpha_1 + \beta_1)^i \omega + (\alpha_1 + \beta_1)^{h-1} (\alpha_1 \epsilon_t^2 + \beta_1 \sigma_t^2). \quad (7.81)$$

Despite similarities to ARCH and GARCH models, forecasts from GJR-GARCH are complicated by the presence of the asymmetric term. If the expected value of the squared shock does not depend on the sign of the return, so that $E[e_t^2 | e_t < 0] = E[e_t^2 | e_t > 0] = 1$, then the probability that $e_{t-1} < 0$ appears in the forecasting formula. When the standardized residuals are normal (or any other symmetric distribution), then this probability is $\frac{1}{2}$. If the density is unknown, this probability must be estimated from the model residuals.

In the GJR-GARCH model, the one-step-ahead forecast is

$$E_t[\sigma_{t+1}^2] = \omega + \alpha_1 \epsilon_t^2 + \alpha_1 \epsilon_t^2 I_{[\epsilon_t < 0]} + \beta_1 \sigma_t^2. \quad (7.82)$$

The two-step-ahead forecast is

$$E_t[\sigma_{t+2}^2] = \omega + \alpha_1 E_t[\epsilon_{t+1}^2] + \alpha_1 E_t[\epsilon_{t+1}^2 I_{\{\epsilon_{t+1} < 0\}}] + \beta_1 E_t[\sigma_{t+1}^2] \quad (7.83)$$

$$= \omega + \alpha_1 E_t[\sigma_{t+1}^2] + \alpha_1 E_t[\epsilon_{t+1}^2 | \epsilon_{t+1} < 0] + \beta_1 E_t[\sigma_{t+1}^2]. \quad (7.84)$$

Assuming the residuals are conditionally normally distributed, then $E_t[\epsilon_{t+1}^2 | \epsilon_{t+1} < 0] = 0.5E[\sigma_{t+1}^2]$.

Multi-step forecasts from other models in the ARCH-family, particularly those that are not linear combinations of ϵ_t^2 , are nontrivial and generally do not have simple recursive formulas. For example, consider forecasting the variance from the simplest nonlinear ARCH-family member, a TARCH(1,0,0) model,

$$\sigma_t = \omega + \alpha_1 |\epsilon_{t-1}| \quad (7.85)$$

As is *always* the case, the 1-step ahead forecast is known at time t ,

$$\begin{aligned} E_t[\sigma_{t+1}^2] &= E_t[(\omega + \alpha_1 |\epsilon_t|)^2] \\ &= E_t[\omega^2 + 2\omega\alpha_1 |\epsilon_t| + \alpha_1^2 \epsilon_t^2] \\ &= \omega^2 + 2\omega\alpha_1 E_t[|\epsilon_t|] + \alpha_1^2 E_t[\epsilon_t^2] \\ &= \omega^2 + 2\omega\alpha_1 |\epsilon_t| + \alpha_1^2 \epsilon_t^2 \end{aligned} \quad (7.86)$$

The 2-step ahead forecast is more complicated and is given by

$$\begin{aligned} E_t[\sigma_{t+2}^2] &= E_t[(\omega + \alpha_1 |\epsilon_{t+1}|)^2] \\ &= E_t[\omega^2 + 2\omega\alpha_1 |\epsilon_{t+1}| + \alpha_1^2 \epsilon_{t+1}^2] \\ &= \omega^2 + 2\omega\alpha_1 E_t[|\epsilon_{t+1}|] + \alpha_1^2 E_t[\epsilon_{t+1}^2] \\ &= \omega^2 + 2\omega\alpha_1 E_t[|e_{t+1}| \sigma_{t+1}] + \alpha_1^2 E_t[e_t^2 \sigma_{t+1}^2] \\ &= \omega^2 + 2\omega\alpha_1 E_t[|e_{t+1}|] E_t[\sigma_{t+1}] + \alpha_1^2 E_t[e_t^2] E_t[\sigma_{t+1}^2] \\ &= \omega^2 + 2\omega\alpha_1 E_t[|e_{t+1}|] (\omega + \alpha_1 |\epsilon_t|) + \alpha_1^2 \cdot 1 \cdot (\omega^2 + 2\omega\alpha_1 |\epsilon_t| + \alpha_1^2 \epsilon_t^2) \end{aligned} \quad (7.87)$$

The challenge in multi-step ahead forecasting of a TARCH model arises since the forecast depends on more than $E_t[e_{t+h}^2] \equiv 1$. In the above example, the forecast depends on both $E_t[e_{t+1}^2] = 1$ and $E_t[|e_{t+1}|]$. When returns are normally distributed, $E_t[|e_{t+1}|] = \sqrt{\frac{2}{\pi}}$, but if the driving innovations have a different distribution, this expectation will differ. The forecast is then, assuming the conditional distribution is normal,

$$E_t[\sigma_{t+2}^2] = \omega^2 + 2\omega\alpha_1 \sqrt{\frac{2}{\pi}} (\omega + \alpha_1 |\epsilon_t|) + \alpha_1^2 (\omega^2 + 2\omega\alpha_1 |\epsilon_t| + \alpha_1^2 \epsilon_t^2). \quad (7.88)$$

The difficulty in multi-step forecasting using “nonlinear” GARCH models – those which involve powers other than two – follows directly from Jensen’s inequality. In the case of TARCH,

$$E_t[\sigma_{t+h}]^2 \neq E_t[\sigma_{t+h}^2] \quad (7.89)$$

or in the general case of an arbitrary power,

$$E_t[\sigma_{t+h}^\delta]^\frac{2}{\delta} \neq E_t[\sigma_{t+h}^2]. \quad (7.90)$$

7.7.1 Evaluating Volatility Forecasts

The evaluation of volatility forecasts is similar to the evaluation of forecasts from conditional mean models with one caveat. In standard time series models, once time $t + h$ has arrived, the value of the variable being forecast is known. However, the value of σ_{t+h}^2 is always unknown in volatility model evaluation and so the realization must be replaced by a proxy. The standard choice is to use the squared return, r_t^2 . This proxy is reasonable if the squared conditional mean is small relative to the variance, a plausible assumption for high-frequency applications to daily or weekly returns. If using longer horizon measurements of returns, e.g., monthly returns, squared residuals ($\hat{\epsilon}_t^2$) estimated from a model for the conditional mean can be used instead. *Realized Variance*, $RV_t^{(m)}$, is an alternative choice to use as a proxy for the unobserved volatility (see section 7.8). Once a choice of proxy has been made, Generalized Mincer-Zarnowitz regressions can be used to assess forecast optimality,

$$r_{t+h}^2 - \hat{\sigma}_{t+h|t}^2 = \gamma_0 + \gamma_1 \hat{\sigma}_{t+h|t}^2 + \gamma_2 z_{1t} + \dots + \gamma_{K+1} z_{Kt} + \eta_t \quad (7.91)$$

where z_{jt} are any instruments known at time t . Common choices for z_{jt} include r_t^2 , $|r_t|$, r_t or indicator variables for the sign of the lagged return. The GMZ regression is testing one key property of a well-specified model: $E_t [r_{t+h}^2 - \hat{\sigma}_{t+h|t}^2] = 0$.

The GMZ regression in equation 7.91 has a heteroskedastic variance, and so a more accurate regression, GMZ-GLS, can be constructed as

$$\frac{r_{t+h}^2 - \hat{\sigma}_{t+h|t}^2}{\hat{\sigma}_{t+h|t}^2} = \gamma_0 \frac{1}{\hat{\sigma}_{t+h|t}^2} + \gamma_1 1 + \gamma_2 \frac{z_{1t}}{\hat{\sigma}_{t+h|t}^2} + \dots + \gamma_{K+1} \frac{z_{Kt}}{\hat{\sigma}_{t+h|t}^2} + \nu_t \quad (7.92)$$

$$\frac{r_{t+h}^2}{\hat{\sigma}_{t+h|t}^2} - 1 = \gamma_0 \frac{1}{\hat{\sigma}_{t+h|t}^2} + \gamma_1 1 + \gamma_2 \frac{z_{1t}}{\hat{\sigma}_{t+h|t}^2} + \dots + \gamma_{K+1} \frac{z_{Kt}}{\hat{\sigma}_{t+h|t}^2} + \nu_t \quad (7.93)$$

by dividing both sides by the time t forecast, $\hat{\sigma}_{t+h|t}^2$ where $\nu_t = \eta_t / \hat{\sigma}_{t+h|t}^2$. Equation 7.93 shows that the GMZ-GLS is a regression of the *generalized error* from a normal likelihood. If one were to use the Realized Variance as the proxy, the GMZ and GMZ-GLS regressions are

$$RV_{t+h} - \hat{\sigma}_{t+h|t}^2 = \gamma_0 + \gamma_1 \hat{\sigma}_{t+h|t}^2 + \gamma_2 z_{1t} + \dots + \gamma_{K+1} z_{Kt} + \eta_t \quad (7.94)$$

and

$$\frac{RV_{t+h} - \hat{\sigma}_{t+h|t}^2}{\hat{\sigma}_{t+h|t}^2} = \gamma_0 \frac{1}{\hat{\sigma}_{t+h|t}^2} + \gamma_1 \frac{\hat{\sigma}_{t+h|t}^2}{\hat{\sigma}_{t+h|t}^2} + \gamma_2 \frac{z_{1t}}{\hat{\sigma}_{t+h|t}^2} + \dots + \gamma_{K+1} \frac{z_{Kt}}{\hat{\sigma}_{t+h|t}^2} + \frac{\eta_t}{\hat{\sigma}_{t+h|t}^2}. \quad (7.95)$$

Diebold-Mariano tests can also be used to test the relative performance of two models. A loss function must be specified when implementing a DM test. Two natural choices for the loss function are MSE,

$$(r_{t+h}^2 - \hat{\sigma}_{t+h|t}^2)^2 \quad (7.96)$$

and QML-loss (which is the *kernel* of the normal log-likelihood),

$$\left(\ln(\hat{\sigma}_{t+h|t}^2) + \frac{r_{t+h}^2}{\hat{\sigma}_{t+h|t}^2} \right). \quad (7.97)$$

The DM statistic is a t-test of the null $H_0 : E[\delta_t] = 0$ where

$$\delta_t = (r_{t+h}^2 - \hat{\sigma}_{A,t+h|t}^2)^2 - (r_{t+h}^2 - \hat{\sigma}_{B,t+h|t}^2)^2 \quad (7.98)$$

in the case of the MSE loss or

$$\delta_t = \left(\ln(\hat{\sigma}_{A,t+h|t}^2) + \frac{r_{t+h}^2}{\hat{\sigma}_{A,t+h|t}^2} \right) - \left(\ln(\hat{\sigma}_{B,t+h|t}^2) + \frac{r_{t+h}^2}{\hat{\sigma}_{B,t+h|t}^2} \right) \quad (7.99)$$

when using QML-loss. Statistically significant positive values of $\bar{\delta} = R^{-1} \sum_{r=1}^R \delta_r$ indicate that B is a better model than A while negative values indicate the opposite (recall R is used to denote the number of out-of-sample observations used to compute the DM statistic). The QML-loss is preferred since it is a “heteroskedasticity corrected” version of the MSE. For more on the evaluation of volatility forecasts using MZ regressions see Patton and Sheppard (2009).

7.8 Realized Variance

Realized Variance (RV) is a new econometric methodology for measuring the variance of asset returns. RV differs from ARCH-models since it does not require a specific model to measure the volatility. Realized Variance instead uses a nonparametric estimator of the variance that is computed *using ultra high-frequency data*.¹⁸

Suppose the log-price process, p_t , is continuously available and is driven by a standard Wiener process with a constant mean and variance,

$$dp_t = \mu dt + \sigma dW_t.$$

The coefficients are normalized so that the return during one day is the difference between p at two consecutive integers (e.g., $p_1 - p_0$ is the first day's return). For the S&P 500 index, $\mu \approx .00031$ and $\sigma \approx .0125$, which correspond to 8% and 20% for the annualized mean and volatility, respectively.

Realized Variance is estimated by sampling p_t throughout the trading day. Suppose that prices on day t were sampled on a regular grid of $m + 1$ points, $0, 1, \dots, m$ and let $p_{i,t}$ denote the i^{th} observation of the log price. The m -sample Realized Variance on day t is defined

$$RV_t^{(m)} = \sum_{i=1}^m (p_{i,t} - p_{i-1,t})^2 = \sum_{i=1}^m r_{i,t}^2. \quad (7.100)$$

Since the price process is a standard Brownian motion, each return is an i.i.d. normal random variable with mean μ/m and variance σ^2/m (or volatility of σ/\sqrt{m}). First, consider the expectation of $RV_t^{(m)}$,

¹⁸Realized Variance was invented somewhere between 1972 and 1997. However, its introduction to modern econometrics clearly dates to the late 1990s (Andersen and Bollerslev, 1998; Andersen, Bollerslev, Diebold, and Labys, 2003; Barndorff-Nielsen and Shephard, 2004).

$$E \left[RV_t^{(m)} \right] = E \left[\sum_{i=1}^m r_{i,t}^2 \right] = E \left[\sum_{i=1}^m \left(\frac{\mu}{m} + \frac{\sigma}{\sqrt{m}} \epsilon_{i,t} \right)^2 \right] \quad (7.101)$$

where $\epsilon_{i,t}$ are i.i.d. standard normal random variables.

$$\begin{aligned} E \left[RV_t^{(m)} \right] &= E \left[\sum_{i=1}^m \left(\frac{\mu}{m} + \frac{\sigma}{\sqrt{m}} \epsilon_{i,t} \right)^2 \right] \\ &= E \left[\sum_{i=1}^m \frac{\mu^2}{m^2} + 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} + \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] \\ &= E \left[\sum_{i=1}^m \frac{\mu^2}{m^2} \right] + E \left[\sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} \right] + E \left[\sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] \\ &= \frac{\mu^2}{m} + \sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} E[\epsilon_{i,t}] + \sum_{i=1}^m \frac{\sigma^2}{m} E[\epsilon_{i,t}^2] \\ &= \frac{\mu^2}{m} + 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \sum_{i=1}^m E[\epsilon_{i,t}] + \frac{\sigma^2}{m} \sum_{i=1}^m E[\epsilon_{i,t}^2] \\ &= \frac{\mu^2}{m} + 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \sum_{i=1}^m 0 + \frac{\sigma^2}{m} \sum_{i=1}^m 1 \\ &= \frac{\mu^2}{m} + \frac{\sigma^2}{m} m \\ &= \frac{\mu^2}{m} + \sigma^2 \end{aligned} \quad (7.102)$$

The expected value is nearly σ^2 , the variance, and it is asymptotically unbiased, $\lim_{m \rightarrow \infty} E \left[RV_t^{(m)} \right] = \sigma^2$. The variance of $RV_t^{(m)}$ can be similarly computed,

$$\begin{aligned} V \left[RV_t^{(m)} \right] &= V \left[\sum_{i=1}^m \frac{\mu^2}{m^2} + 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} + \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] \\ &= V \left[\sum_{i=1}^m \frac{\mu^2}{m^2} \right] + V \left[\sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} \right] + V \left[\sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] + 2 \text{Cov} \left[\sum_{i=1}^m \frac{\mu^2}{m^2}, \sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} \right] \\ &\quad + 2 \text{Cov} \left[\sum_{i=1}^m \frac{\mu^2}{m^2}, \sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] + 2 \text{Cov} \left[\sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t}, \sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right]. \end{aligned} \quad (7.103)$$

First, the variance and covariance terms that involve the mean term are all zero,

$$V \left[\sum_{i=1}^m \frac{\mu^2}{m^2} \right] = \text{Cov} \left[\sum_{i=1}^m \frac{\mu^2}{m^2}, \sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} \right] = \text{Cov} \left[\sum_{i=1}^m \frac{\mu^2}{m^2}, \sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] = 0,$$

since $\frac{\mu^2}{m^2}$ is a constant. The remaining covariance term also has expectation 0 since $\epsilon_{i,t}$ are i.i.d. standard normal and so have a skewness of 0,

$$\text{Cov} \left[\sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t}, \sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] = 0$$

The other two terms can be shown to be (left as exercises)

$$\begin{aligned} V \left[\sum_{i=1}^m 2 \frac{\mu\sigma}{m^{\frac{3}{2}}} \epsilon_{i,t} \right] &= 4 \frac{\mu^2 \sigma^2}{m^2} \\ V \left[\sum_{i=1}^m \frac{\sigma^2}{m} \epsilon_{i,t}^2 \right] &= 2 \frac{\sigma^4}{m} \end{aligned}$$

and so

$$V \left[RV_t^{(m)} \right] = 4 \frac{\mu^2 \sigma^2}{m^2} + 2 \frac{\sigma^4}{m}. \quad (7.104)$$

The variance is decreasing as $m \rightarrow \infty$, $RV_t^{(m)}$ is asymptotically unbiased, and so $RV_t^{(m)}$ is a consistent estimator of σ^2 .

In the empirically realistic case where the price process has a time-varying drift and stochastic volatility,

$$dp_t = \mu_t dt + \sigma_t dW_t,$$

$RV_t^{(m)}$ is a consistent estimator of the *integrated variance*,

$$\lim_{m \rightarrow \infty} RV_t^{(m)} \xrightarrow{P} \int_t^{t+1} \sigma_s^2 ds. \quad (7.105)$$

The integrated variance measures the average variance of the measurement interval, usually a day.

If the price process contains jumps, $RV_t^{(m)}$ is still a consistent estimator although its limit is the *quadratic variation* rather than the integrated variance, and so

$$\lim_{m \rightarrow \infty} RV_t^{(m)} \xrightarrow{P} \int_t^{t+1} \sigma_s^2 ds + \sum_{t \leq 1} \Delta J_s^2. \quad (7.106)$$

where $\sum_{t \leq 1} \Delta J_s^2$ is the sum of the squared jumps if any. Similar results hold if the price process exhibits leverage (instantaneous correlation between the price and the variance). The two conditions for $RV_t^{(m)}$ to be a reasonable method to estimate the integrated variance on day t are essentially that the price process, p_t , is arbitrage-free and that the efficient price is observable. Empirical evidence suggests that prices of liquid asset are compatible with the first condition. The second condition is violated since assets trade at either the best bid or best ask price – neither of which is the efficient price.

7.8.1 Implementing Realized Variance

In practice, naïve implementations of Realized Variance do not perform well. The most pronounced challenge is that observed prices are contaminated by noise; there is no single price, and traded prices are only observed at the bid and the ask. This feature of asset price transactions produces bid-ask bounce where consecutive prices oscillate between the two. Consider a simple model of bid-ask bounce where returns are computed as the log difference in observed prices composed of the true (unobserved) efficient prices, $p_{i,t}^*$, contaminated by an independent mean zero shock, $v_{i,t}$,

$$p_{i,t} = p_{i,t}^* + v_{i,t}.$$

The shock $v_{i,t}$ captures the difference between the efficient price and the observed prices which are always on the bid or ask price.

The i^{th} observed return, $r_{i,t}$ can be decomposed into the actual (unobserved) return $r_{i,t}^*$ and an independent noise term $\eta_{i,t} = v_{i,t} - v_{i-1,t}$,

$$\begin{aligned} p_{i,t} - p_{i-1,t} &= (p_{i,t}^* + v_{i,t}) - (p_{i-1,t}^* + v_{i-1,t}) \\ p_{i,t} - p_{i-1,t} &= (p_{i,t}^* - p_{i-1,t}^*) + (v_{i,t} - v_{i-1,t}) \\ r_{i,t} &= r_{i,t}^* + \eta_{i,t} \end{aligned} \tag{7.107}$$

The error in the observed return process, $\eta_{i,t} = v_{i,t} - v_{i-1,t}$, is a MA(1) and so is serially correlated.

Computing the RV from returns contaminated by noise has an unambiguous effect on Realized Variance; RV is biased upward.

$$\begin{aligned} RV_t^{(m)} &= \sum_{i=1}^m r_{i,t}^2 \\ &= \sum_{i=1}^m (r_{i,t}^* + \eta_{i,t})^2 \\ &= \sum_{i=1}^m r_{i,t}^{*2} + 2r_{i,t}^* \eta_{i,t} + \eta_{i,t}^2 \\ &\approx \widehat{RV}_t + m\tau^2 \end{aligned} \tag{7.108}$$

where τ^2 is the variance of $\eta_{i,t}$ and \widehat{RV}_t is the Realized Variance that would be computed if the efficient returns could be observed. The bias is increasing in the number of samples (m) and can be substantial for assets with large bid-ask spreads.

The simplest “solution” to the bias is to avoid the issue using *sparse sampling*, i.e., not using all of the observed prices. The noise imposes limits on m to ensure that the bias is small relative to the integrated variance. In practice the maximum m is always much higher than 1 – a single open-to-close return – and is typically somewhere between 13 (30-minute returns on a stock listed on the NYSE) and 390 (1-minute returns), and so even when $RV_t^{(m)}$ is not consistent, it is still a better proxy, often substantially, for the latent variance on day t than r_t^2 (the “1-sample Realized Variance”, see Bandi and Russell (2008)). The signal-to-noise ratio (which measures the

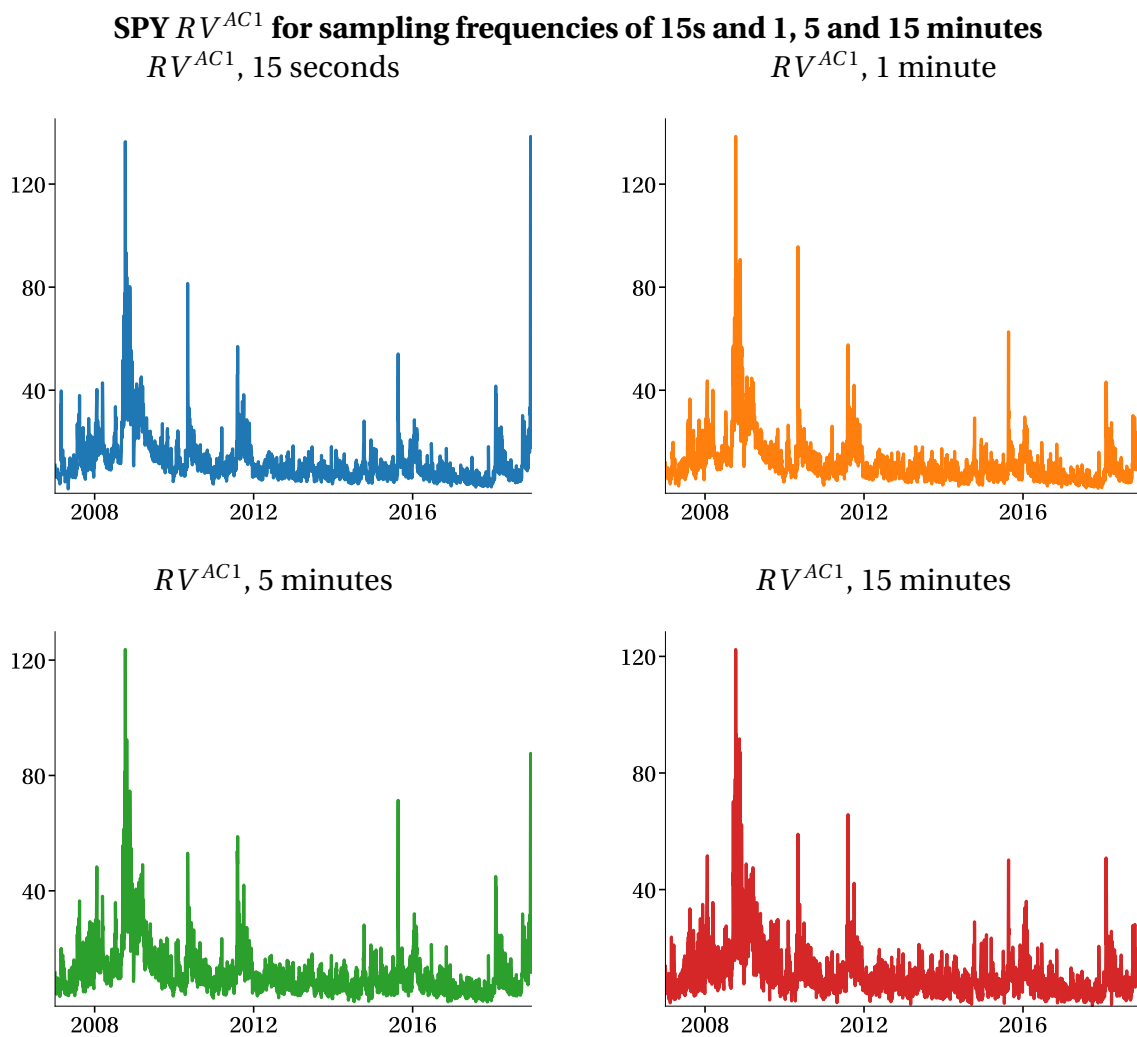


Figure 7.9: The four panels of this figure contain a noise-robust version Realized Variance, RV^{AC1} , for every day the market was open from January 2007 until December 2018 transformed into annualized volatility. The 15-second RV^{AC1} is better behaved than the 15-second RV .

ratio of useful information to pure noise) is approximately 1 for RV but is between .05 and .1 for r_t^2 . In other words, RV is 10-20 times more precise than squared daily returns (Andersen and Bollerslev, 1998).

Another simple and effective method is to filter the data using an MA(1). Transaction data contain a strong negative MA due to bid-ask bounce, and so RV computed using the errors ($\hat{\epsilon}_{i,t}$) from a model,

$$r_{i,t} = \theta \epsilon_{i-1,t} + \epsilon_{i,t} \quad (7.109)$$

eliminates much of the bias. A better method to remove the bias is to use an estimator known as RV^{AC1} which is similar to a Newey-West estimator.




Figure 7.10: The volatility signature plot for the RV shows a clear trend. Based on visual inspection, it would be difficult to justify sampling more frequently than 30 seconds. Unlike the volatility signature plot of the RV , the signature plot of RV^{AC1} does not monotonically increase with the sampling frequency except when sampling every second, and the range of the values is considerably smaller than in the RV signature plot.

$$RV_t^{AC1(m)} = \sum_{i=1}^m r_{i,t}^2 + 2 \sum_{i=2}^m r_{i,t} r_{i-1,t} \quad (7.110)$$

In the case of a constant drift, constant volatility Brownian motion subject to bid-ask bounce, this estimator can be shown to be unbiased, although it is not consistent in large samples. A more general class of estimators that use a kernel structure that can be tuned to match the characteristics of specific asset prices and which are consistent as $m \rightarrow \infty$ even in the presence of noise has been introduced in Barndorff-Nielsen, Hansen, Lunde, and Shephard (2008).¹⁹

Another problem for Realized Variance is that prices are not available at regular intervals.

¹⁹The Newey-West estimator is a particular implementation of a broad class of estimators known as kernel variance estimators. They all share the property that they are weighted sums of autocovariances where a kernel function determines the weights.

Fortunately, this issue has a simple solution: *last price interpolation*. Last price interpolation sets the price at time t to the last observed price p_τ where τ is the largest time index less where $\tau \leq t$. Other interpolation schemes produce bias in RV . Consider, for example, linear interpolation which sets prices at time- t price to $p_t = wp_{\tau_1} + (1 - w)p_{\tau_2}$ where τ_1 is the time subscript of the last observed price before t and τ_2 is the time subscript of the first price after time t , and the interpolation weight is $w = (\tau_2 - t)/(\tau_2 - \tau_1)$. The averaging of prices in linear interpolation – which effectively produces a smoother price path than the efficient price path – produces a notable downward bias in RV .

Finally, most markets do not operate 24 hours a day, and RV cannot be computed when markets are closed. The standard procedure is to augment high-frequency returns with the squared close-to-open return to construct an estimate of the total variance. The close-to-close (CtC) RV is then defined

$$RV_{\text{CtC},t}^{(m)} = r_{\text{CtO},t}^2 + RV_t^{(m)} \quad (7.111)$$

where $r_{\text{CtO},t}^2$ is the return between the close on day $t - 1$ and the market open on day t . Since the overnight return is not measured frequently, the adjusted RV must be treated as a random variable (and not an observable). An improved method to handle the overnight return has been proposed in Hansen and Lunde (2005) and Hansen and Lunde (2006) which weighs the overnight squared return by λ_1 and the daily Realized Variance by λ_2 to produce an estimator with a lower mean-square error,

$$\widetilde{RV}_{\text{CtC},t}^{(m)} = \lambda_1 r_{\text{CtO},t}^2 + \lambda_2 RV_t^{(m)}.$$

7.8.2 Modeling RV

If RV is observable, then it can be modeled using standard time series tools such as ARMA models. This approach has been widely used in the academic literature although there are issues in treating the RV “as-if” it is the variance. If RV has measurement error, then parameter estimates in ARMA models suffer from an errors-in-variables problem, and the estimated coefficient are biased (see chapter 4). Corsi (2009) proposed the *heterogeneous autoregression* (HAR) as a simple method to capture the dynamics in RV in a parsimonious model. The standard HAR models the RV as a function of the RV in the previous day, the average RV over the previous week, and the average RV over the previous month (22 days). The HAR in levels is then

$$RV_t = \phi_0 + \phi_1 RV_{t-1} + \phi_5 \overline{RV}_{t-5} + \phi_2 2\overline{RV}_{t-22} + \epsilon_t \quad (7.112)$$

where $\overline{RV}_{t-5} = \frac{1}{5} \sum_{i=1}^5 RV_{t-i}$ and $\overline{RV}_{t-22} = \frac{1}{22} \sum_{i=1}^{22} 2RV_{t-i}$ (suppressing the (m) terms). The HAR is also commonly estimated in logs,

$$\ln RV_t = \phi_0 + \phi_1 \ln RV_{t-1} + \phi_5 \ln \overline{RV}_{t-5} + \phi_2 2 \ln \overline{RV}_{t-22} + \epsilon_t. \quad (7.113)$$

HARs are technically AR(22) models with many parameter restrictions. These restrictions maintain parsimony while allowing HARs to capture both the high degree of persistence in volatility (through the 22-day moving average) and short term dynamics (through the 1-day and 5-day terms).

The alternative is to model RV using ARCH-family models, which can be interpreted as multiplicative error models for *any* non-negative process, not only squared returns (Engle, 2002a).²⁰ Standard statistical software can be used to model RV as an ARCH process by defining $\tilde{r}_t = \text{sign}(r_t)\sqrt{RV_t}$ where $\text{sign}(r_t)$ is 1 if the end-of-day return is positive or -1 otherwise. The transformed RV , \tilde{r}_t , is the signed square root of the Realized Variance on day t . Any ARCH-family model can be applied to these transformed values. For example, when modeling the variance evolution as a GJR-GARCH(1,1,1) process,

$$\sigma_t^2 = \omega + \alpha_1 \tilde{r}_{t-1}^2 + \gamma_1 \tilde{r}_{t-1}^2 I_{[\tilde{r}_{t-1} < 0]} + \beta_1 \sigma_{t-1}^2 \quad (7.114)$$

which is equivalently expressed in terms of Realized Variance as

$$\sigma_t^2 = \omega + \alpha_1 RV_{t-1} + \gamma_1 RV_{t-1} I_{[r_{t-1} < 0]} + \beta_1 \sigma_{t-1}^2. \quad (7.115)$$

Maximum likelihood estimation, assuming normally distributed errors, can be used to estimate the parameters of this model. This procedure solves the errors-in-variables problem present when RV is treated as observable and facilitates modeling RV using standard software. Inference and the method to build a model are unaffected by the change from end-of-day returns to the transformed RV .

7.8.3 Realized Variance of the S&P 500

Returns on S&P 500 Depository Receipts, known as SPiDeRs (NYSEARCA:SPY) is used to illustrate the gains and pitfalls of RV . Price data was taken from TAQ and includes every transaction between January 2007 until December 2018, a total of 3,020 days. SPDRs track the S&P 500 and are among the most liquid assets in the U.S. market with an average volume of 150 million shares per day. There were more than 100,000 trades on a typical day throughout the sample, which is more than 4 per second. TAQ data contain errors, and observations were filtered by removing the prices outside the daily high or low from an audited database. Only trade prices that occurred during the usual trading hours of 9:30 – 16:00 were retained.

The primary tool for examining different Realized Variance estimators is the volatility signature plot.

Definition 7.10 (Volatility Signature Plot). The volatility signature plot displays the time-series average of Realized Variance

$$\overline{RV}_t^{(m)} = T^{-1} \sum_{t=1}^T RV_t^{(m)}$$

as a function of the number of samples, m . An equivalent representation displays the amount of time, whether in calendar time or tick time (number of trades between observations) along the X-axis.

Figures 7.11 and 7.9 contain plots of the *annualized volatility* constructed from the RV and RV^{AC1} . The estimates have been annualized to facilitate interpretation. Figures 7.11 shows that the 15-second RV is larger than the RV sampled at 1, 5 or 15 minutes and that the 1 and 5 minute

²⁰ARCH-family models have, for example, been successfully applied to both durations (time between trades) and hazards (number of trades in an interval of time), two non-negative processes.

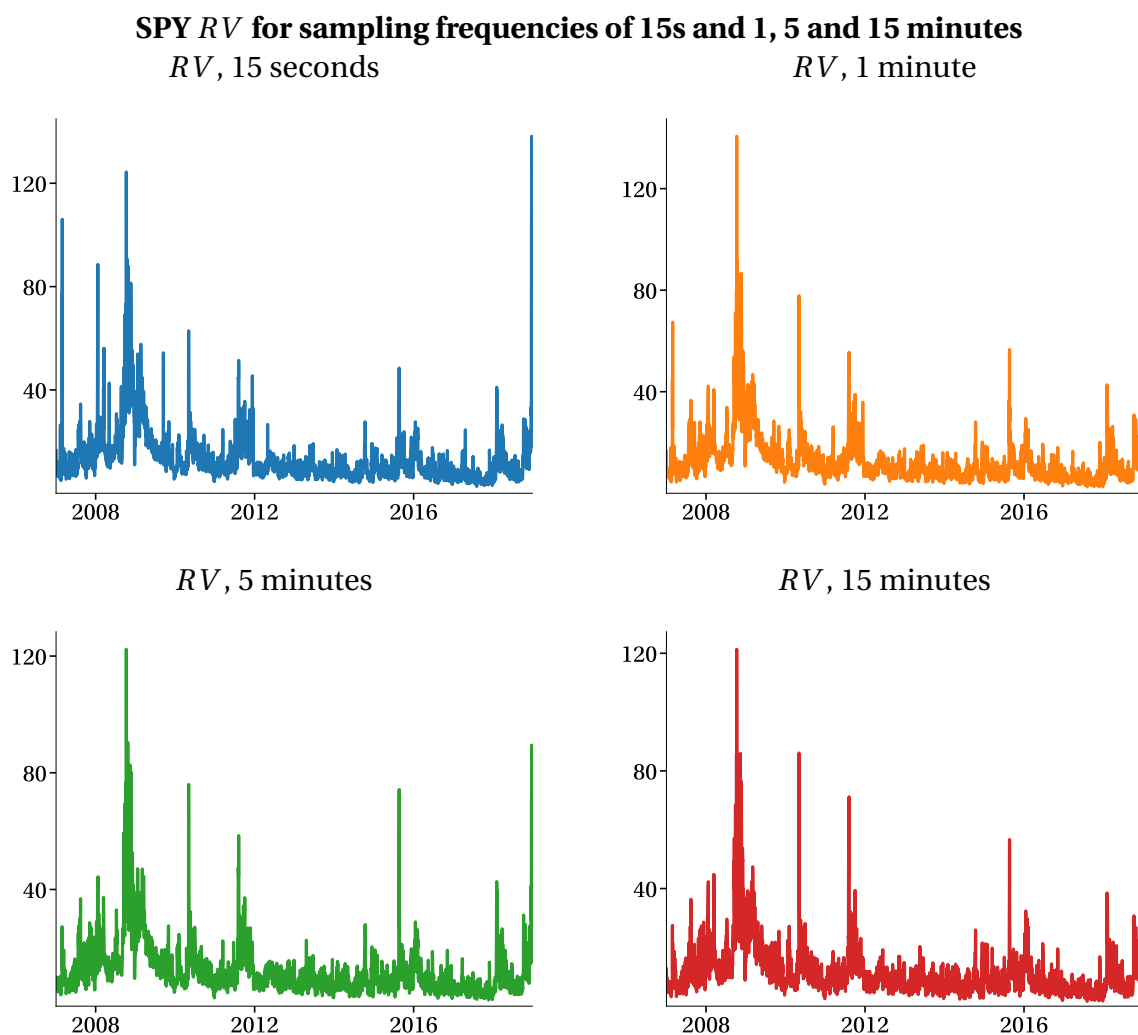


Figure 7.11: The four panels of this figure contain the Realized Variance for every day the market was open from January 2007 until December 2018. The estimated RV have been transformed into annualized volatility ($\sqrt{252 \cdot RV_t^{(m)}}$). While these plots appear superficially similar, the 1- and 5-minute RV are the most precise and the 15-second RV is biased upward.

RV are less noisy than the 15-minute RV . These plots provide some evidence that sampling more frequently than 15 minutes may be desirable. The two figures show that there is a reduction in the scale of the 15-second RV^{AC1} relative to the 15-second RV . The 15-second RV is heavily influenced by the noise in the data (bid-ask bounce) while the RV^{AC1} is less affected.

Figures 7.10 and 7.10 contain the annualized volatility signature plot for RV and RV^{AC1} , respectively. The dashed horizontal line depicts the volatility computed using the standard variance estimator computed from open-to-close returns. There is a striking difference between the two figures. The RV volatility signature plot diverges when sampling more frequently than 30 seconds while the RV^{AC1} plot is flat except at the highest sample frequency. RV^{AC1} appears to allow sampling every 5 seconds – 6 times more frequently than RV . This is a common finding

when comparing RV^{AC1} to RV across a wide range of asset price data.

7.9 Implied Volatility and VIX

Implied volatility differs from other measures in that it is both market-based and forward-looking. Implied volatility was originally conceived as the “solution” to the Black-Scholes options pricing formula where all values except the volatility are observable. Recall that the Black-Scholes formula is derived from an assuming that stock prices follow a geometric Brownian motion plus drift,

$$dS_t = \mu S_t dt + \sigma S_t dW_t \quad (7.116)$$

where S_t is the time t stock prices, μ is the drift, σ is the (constant) volatility, and dW_t is a Wiener process. Under some additional assumptions sufficient to ensure no arbitrage, the price of a call option can be shown to be

$$\begin{aligned} C_t(T, K) &= S_t \Phi(d_1) + K e^{-rT} \Phi(d_2) \\ d_1 &= \frac{\ln(S_t/K) + (r + \sigma^2/2) T}{\sigma \sqrt{T}} \\ d_2 &= \frac{\ln(S_t/K) + (r - \sigma^2/2) T}{\sigma \sqrt{T}} \end{aligned} \quad (7.117)$$

where K is the strike price, T is the time to maturity, reported in years, r is the risk-free interest rate, and $\Phi(\cdot)$ is the normal CDF. The price of a call option is monotonic in the volatility, and so the formula can be inverted to express the volatility as a function of the call price and other observables. The implied volatility,

$$\sigma_t^{\text{Implied}} = g(C_t(T, K), S_t, K, T, r), \quad (7.118)$$

is the expected volatility between t and T under the risk-neutral measure (which is the same as under the physical when volatility is constant).²¹

7.9.1 The smile

When computing the Black-Scholes implied volatility across a range of strikes, the volatility usually resembles a “smile” (higher IV for out-of-the-money options than in the money) or “smirk” (higher IV for out-of-the-money puts). This pattern emerges since asset returns are heavy-tailed (“smile”) and skewed (“smirk”). The BSIV is derived under an assumption that the asset price follows a *geometric Brownian motion* so that the log returns are assumed to be normal. The smile reflects misspecification of the model underlying the Black-Scholes option pricing formula. Figure 7.12 shows the smile in the BSIV for SPY out-of-the-month options on January 15, 2017. The

²¹The implied volatility is computed by numerically inverting the B-S pricing formula, or using some other approximation..




Figure 7.12: Plot of the Black-Scholes implied volatility “smile” on January 15, 2018, based on options on SPY expiring on February 2, 2018.

x-axis rescaled from the strike price to *moneyness* by dividing the strike by the spot price. The current spot price is 100, smaller values indicate strikes below the current price (out-of-the-money puts), and positive values are strikes above the current price (out-of-the-money calls).

7.9.2 Model-Free Volatility

B-S implied volatility suffers from three key issues:

- **Derived under constant volatility:** The returns on most asset prices exhibit conditional heteroskedasticity, and time-variation in the volatility of returns generates heavy tails which increases the probability of a large asset price change.
- **Leverage effects are ruled out:** Leverage, or negative correlation between the price and volatility of an asset, can generate negative skewness. This feature of asset prices increases the probability of extreme negative returns relative to the log-normal price process assumed in the B-S option pricing formula.
- **No jumps:** Jumps are also an empirical fact of most asset prices. Jumps, like time-varying volatility, increase the chance of seeing an extreme return.

The consequences of these limits are that, contrary to what the model underlying the B-S implies, B-S implied volatilities are not constant across strike prices, and so cannot be interpreted as market-based estimated of volatility.

Model-free implied volatility (MFIV) was been developed as an alternative to B-S implied volatility by Demeterfi et al. (1999) and Britten-Jones and Neuberger (2000) with an important extension to jump processes and practical implementation details provided by Jiang and Tian (2005). These estimators build on Breeden and Litzenberger (1978) which contains key result that demonstrates how option prices are related to the risk-neutral measure – the distribution of asset price returns after removing risk premia. Suppose that the risk-neutral measure \mathbb{Q} exists and is unique. Then, under the risk-neutral measure, it must be the case that

$$\frac{\partial S_t}{S_t} = \sigma(t, \cdot) dW_t \quad (7.119)$$

is a martingale where $\sigma(t, \cdot)$ is a (possibly) time-varying volatility process that may depend on the stock price or other state variables. From the relationship, the price of a call option can be computed as

$$C(t, K) = E_{\mathbb{Q}} [(S_t - K)^+] \quad (7.120)$$

for $t > 0$, $K > 0$ where the function $(x)^+ = \max(x, 0)$. Thus

$$C(t, K) = \int_K^{\infty} (S_t - K) \phi_t(S_t) dS_t \quad (7.121)$$

where $\phi_t(\cdot)$ is the risk-neutral measure. Differentiating with respect to K ,

$$\frac{\partial C(t, K)}{\partial K} = - \int_K^{\infty} \phi_t(S_t) dS_t. \quad (7.122)$$

Differentiating this expression again with respect to K (note K in the lower integral bound),

$$\frac{\partial^2 C(t, K)}{\partial K^2} = \phi_t(K), \quad (7.123)$$

and so that the risk-neutral density can be recovered from options prices. This result provides a basis for nonparametrically estimating the risk-neutral density from observed options prices (see, e.g., Aït-Sahalia and Lo (1998)). Another consequence of this result is that the expected (under \mathbb{Q}) variation in a stock price over the interval $[t_1, t_2]$ measure can be recovered from

$$E_{\mathbb{Q}} \left[\int_{t_1}^{t_2} \left(\frac{\partial S_t}{S_t} \right)^2 \right] = 2 \int_0^{\infty} \frac{C(t_2, K) - C(t_1, K)}{K^2} dK. \quad (7.124)$$

This expression cannot be directly implemented to recover the expected volatility since it requires a continuum of strike prices.

Equation 7.124 assumes that the risk-free rate is 0. When it is not, a similar result can be derived using the forward price

$$E_{\mathbb{F}} \left[\int_{t_1}^{t_2} \left(\frac{\partial F_t}{F_t} \right)^2 \right] = 2 \int_0^{\infty} \frac{C^F(t_2, K) - C^F(t_1, K)}{K^2} dK \quad (7.125)$$

where \mathbb{F} is the forward probability measure – that is, the probability measure where the forward price is a martingale and $C^F(\cdot, \cdot)$ is used to denote that this option is defined on the forward price. Additionally, when t_1 is 0, as is usually the case, the expression simplifies to

$$E_{\mathbb{F}} \left[\int_0^t \left(\frac{\partial F_t}{F_t} \right)^2 \right] = 2 \int_0^\infty \frac{C^F(t, K) - (F_0 - K)^+}{K^2} dK. \quad (7.126)$$

A number of important caveats are needed for employing this relationship to compute MFIV from option prices:

- Spot rather than forward prices. Because spot prices are usually used rather than forwards, the dependent variable needs to be redefined. If interest rates are non-stochastic, then define $B(0, T)$ to be the price of a bond today that pays \$1 time T . Thus, $F_0 = S_0/B(0, T)$, is the forward price and $C^F(T, K) = C(T, K)/B(0, T)$ is the forward option price. With the assumption of non-stochastic interest rates, the model-free implied volatility can be expressed

$$E_{\mathbb{F}} \left[\int_0^t \left(\frac{\partial S_t}{S_t} \right)^2 \right] = 2 \int_0^\infty \frac{C(t, K)/B(0, T) - (S_0/B(0, T) - K)^+}{K^2} dK \quad (7.127)$$

or equivalently using a change of variables as

$$E_{\mathbb{F}} \left[\int_0^t \left(\frac{\partial S_t}{S_t} \right)^2 \right] = 2 \int_0^\infty \frac{C(t, K/B(0, T)) - (S_0 - K)^+}{K^2} dK. \quad (7.128)$$

- Discretization. Because only finitely many options prices are available, the integral must be approximated using a discrete grid. Thus the approximation

$$E_{\mathbb{F}} \left[\int_0^t \left(\frac{\partial S_t}{S_t} \right)^2 \right] = 2 \int_0^\infty \frac{C(t, K/B(0, T)) - (S_0 - K)^+}{K^2} dK \quad (7.129)$$

$$\approx \sum_{m=1}^M [g(T, K_m) + g(T, K_{m-1})] (K_m - K_{m-1}) \quad (7.130)$$

The is used where

$$g(T, K) = \frac{C(t, K/B(0, T)) - (S_0 - K)^+}{K^2} \quad (7.131)$$

If the option tree is rich, this should not pose a significant issue. For option trees on individual firms, asset-specific study (for example, using data-calibrated Monte Carlo experiment) may be needed to ascertain whether the MFIV is a good estimate of the volatility under the forward measure.

- Maximum and minimum strike prices. The integral cannot be implemented from 0 to ∞ , and so the implied volatility has a downward bias due to the effect of the tails. In rich options trees, such as for the S&P 500, this issue is minor.

7.9.3 VIX

The VIX – Volatility Index – is a volatility measure produced by the Chicago Board Options Exchange (CBOE). It is computed using a “model-free” like estimator which uses both call and put prices.²² The VIX is an estimator of the price of a variance swap, which applies put-call parity to the previous expression to produce

$$\frac{2}{T} \exp(rT) \left(\int_0^{F_0} \frac{P(t, K/B(0, T))}{K^2} dK + \int_{F_0}^{\infty} \frac{C(t, K/B(0, T))}{K^2} dK \right).$$

The term $(S_0 - K)^+$ drops out of this expression since it only used out-of-the-money options.

The VIX is computed according to

$$\sigma^2 = \frac{2}{T} \exp(rT) \sum_{i=1}^N \frac{\Delta K_i}{K_i^2} Q(K_i) - \frac{1}{T} \left(\frac{F_0}{K_0} - 1 \right)^2 \quad (7.132)$$

where T is the time to expiration of the options used, F_0 is the forward price which is computed from index option prices, K_i is the strike of the i^{th} out-of-the-money option, $\Delta K_i = (K_{i+1} - K_{i-1})/2$ is half of the distance of the interval surrounding the option with a strike price of K_i , K_0 is the strike of the option immediately below the forward level, F_0 , r is the risk-free rate and $Q(K_i)$ is the mid-point of the bid and ask for the call or put used at strike K_i . The forward index price is extracted using put-call parity as $F_0 = K_0 + \exp(rT)(C_0 - P_0)$ where K_0 is the strike price where the price difference between put and call is smallest, and C_0 and P_0 are, respectively, the call and put prices at this node. The VIX is typically calculated from options at the two maturities closes to the 30-day horizon (for example 28- and 35-days when using options that expire weekly). More details on the implementation of the VIX can be found in the CBOE whitepaper (CBOE, 2003).

The first term in the formula for the VIX can be viewed as

$$\frac{\Delta K_i}{K_i^2} Q(K_i) = \underbrace{\frac{\Delta K_i}{K_i}}_{\% \text{ width of interval}} \times \underbrace{\frac{Q(K_i)}{K_i}}_{\% \text{ option premium}},$$

so that the implied variance depends on only the option premium as a percent of the strike price. The division in the second term by K_0 similarly transforms the forward price to a percentage of strike measure. Each of these terms is width time height (premium), and so the VIX is the area below the out-of-the-money option pricing curve. When volatility is higher, all options are more valuable, and so there is more area below the curve. Figure 7.13 illustrates this area using option prices computed from the Black-Scholes formula for volatilities of 20% and 60%.

7.9.4 Computing the VIX from Black-Scholes prices

Put and call options values were computed from the Black-Scholes option pricing formula for an underlying with a price of \$100, an option time to maturity of a month ($T = 1/12$), a volatility of 20%, and a risk-free rate of 2%. Figure 7.13 plots the put and call options values from the Black-Scholes formula. The solid lines indicate the options that are out-of-the-money – puts

²²The VIX is based exclusively on out-of-the-money prices, so calls are used for strikes above the current price and puts are used for strikes below the current price.




Figure 7.13: Option prices generated from the Black-Scholes pricing formula for an underlying with a price of \$100 with a volatility of 20% or 60% (bottom). The options expire in 1 month ($T = 1/12$), and the risk-free rate is 2%. The solid lines show the out-of-the-money options that are used to compute the VIX. The solid markers show the values where the option price to be at least \$0.01 using a \$4 grid of strike prices.

with strike prices below \$100 or calls with strikes above \$100 – that are used to compute the VIX. The dotted lines show the option prices that are in-the-money. The values in Table 7.7 show all strikes where the out-of-the-month option price was at least \$0.01. These values are marked in Figure 7.13. The VIX is computed using the out-of-the-money option price $Q(K_i)$ rescaled by $2/T \exp(rT) \Delta K_i / K_i^2 = 2/1/12 \exp(.02/12) \times 4/K_i$ since the strikes are measured every \$4. The final line shows the total – 0.0430. The VIX index computed from these values is then $100 \times \sqrt{0.0430 - 3.338 \times 10^{-5}} = 20.75\%$, which is close to the true value of 20%. The second term in the square root is the adjustment $1/T (F/K_0 - 1)^2$ which is small. The small difference between the MFIV and the true volatility of 20% is due to discretization error since the strikes are only observed every \$4 and truncation error since only options with values larger than \$0.01 were used. The bottom panel of Figure 7.13 plots the option prices and highlights the area estimated by the VIX formula when the asset price volatility is 60%.

Strike	Call	Put	Abs. Diff.	VIX Contrib.
88	12.17	0.02	12.15	0.0002483
92	8.33	0.17	8.15	0.0019314
96	4.92	0.76	4.16	0.0079299
100	2.39	2.22	0.17	0.0221168
104	0.91	4.74	3.83	0.0080904
108	0.27	8.09	7.82	0.0022259
112	0.06	11.88	11.81	0.0004599
116	0.01	15.82	15.81	7.146e-05
Total				0.0430742

Table 7.7: Option prices generated from the Black-Scholes pricing formula for an underlying with a price of \$100 with a volatility of 20%. The options expire in 1 month ($T = 1/12$), and the risk-free rate is 2%. The third column shows the absolute difference which is used to determine K_0 in the VIX formula. The final column contains the contribution of each option to the VIX as measured by $^{2/T} \exp(rT) \Delta K_i / K_i^2 \times Q(K_i)$.

7.9.5 Empirical Relationships

The daily VIX series from January 1990 until December 2018 is plotted in Figure 7.14 against a 22-day *forward* moving average computed as

$$\sigma_t^{MA} = \sqrt{\frac{252}{22} \sum_{i=0}^{21} r_{t+i}^2}.$$

The second panel shows the difference between these two series. The VIX is consistently, but not uniformly, higher than the forward volatility. This relationship highlights both a feature and a drawback of using a measure of the volatility computed under the risk-neutral measure: it captures a (possibly) time-varying risk premium. This risk premium captures investor compensation for changes in volatility (volatility of volatility) and jump risks.

7.A Kurtosis of an ARCH(1)

The necessary steps to derive the kurtosis of an ARCH(1) process are

$$\begin{aligned}
E[\epsilon_t^4] &= E[E_{t-1}[\epsilon_t^4]] \\
&= E[3(\omega + \alpha_1 \epsilon_{t-1}^2)^2] \\
&= 3E[(\omega + \alpha_1 \epsilon_{t-1}^2)^2] \\
&= 3E[\omega^2 + 2\omega\alpha_1 \epsilon_{t-1}^2 + \alpha_1^2 \epsilon_{t-1}^4] \\
&= 3(\omega^2 + \omega\alpha_1 E[\epsilon_{t-1}^2] + \alpha_1^2 E[\epsilon_{t-1}^4]) \\
&= 3\omega^2 + 6\omega\alpha_1 E[\epsilon_{t-1}^2] + 3\alpha_1^2 E[\epsilon_{t-1}^4].
\end{aligned} \tag{7.133}$$

Using μ_4 to represent the expectation of the fourth power of ϵ_t ($\mu_4 = E[\epsilon_t^4]$),

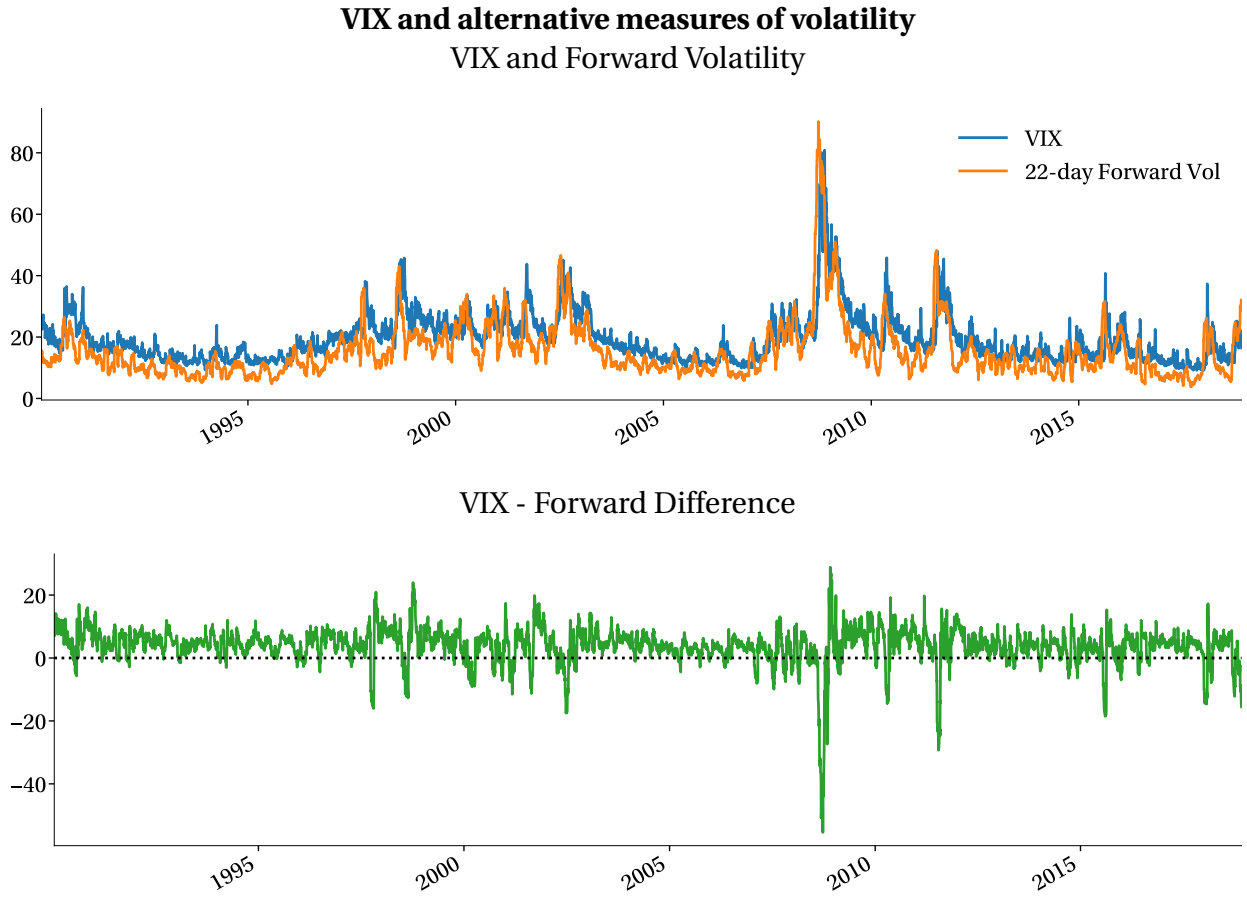


Figure 7.14: Plots of the VIX against a TARCH-based estimate of the volatility (top panel) and a 22-day forward moving average (bottom panel). The VIX is consistently above both measures reflecting the presence of a risk premium that compensates for time-varying volatility and jumps in the market return.

$$\begin{aligned}
 E[\epsilon_t^4] - 3\alpha_1^2 E[\epsilon_{t-1}^4] &= 3\omega^2 + 6\omega\alpha_1 E[\epsilon_{t-1}^2] \\
 \mu_4 - 3\alpha_1^2 \mu_4 &= 3\omega^2 + 6\omega\alpha_1 \bar{\sigma}^2 \\
 \mu_4(1 - 3\alpha_1^2) &= 3\omega^2 + 6\omega\alpha_1 \bar{\sigma}^2 \\
 \mu_4 &= \frac{3\omega^2 + 6\omega\alpha_1 \bar{\sigma}^2}{1 - 3\alpha_1^2} \\
 \mu_4 &= \frac{3\omega^2 + 6\omega\alpha_1 \frac{\omega}{1-\alpha_1}}{1 - 3\alpha_1^2} \\
 \mu_4 &= \frac{3\omega^2(1 + 2\frac{\alpha_1}{1-\alpha_1})}{1 - 3\alpha_1^2} \\
 \mu_4 &= \frac{3\omega^2(1 + \alpha_1)}{(1 - 3\alpha_1^2)(1 - \alpha_1)}.
 \end{aligned} \tag{7.134}$$

This derivation makes use of the same principals as the intuitive proof and the identity that $\bar{\sigma}^2 = \omega/(1 - \alpha_1)$. The final form highlights two important issues: first, μ_4 (and thus the kurtosis) is only finite if $1 - 3\alpha_1^2 > 0$ which requires that $\alpha_1 < \sqrt{\frac{1}{3}} \approx .577$, and second, the kurtosis, $\kappa = \frac{E[\epsilon_t^4]}{E[\epsilon_t^2]^2} = \frac{\mu_4}{\bar{\sigma}^2}$, is always greater than 3 since

$$\begin{aligned} \kappa &= \frac{E[\epsilon_t^4]}{E[\epsilon_t^2]^2} \\ &= \frac{\frac{3\omega^2(1+\alpha_1)}{(1-3\alpha_1^2)(1-\alpha_1)}}{\frac{\omega^2}{(1-\alpha_1)^2}} \\ &= \frac{3(1-\alpha_1)(1+\alpha_1)}{(1-3\alpha_1^2)} \\ &= \frac{3(1-\alpha_1^2)}{(1-3\alpha_1^2)} > 3. \end{aligned} \tag{7.135}$$

Finally, the variance of ϵ_t^2 can be computed noting that for any variable Y , $V[Y] = E[Y^2] - E[Y]^2$, and so

$$\begin{aligned} V[\epsilon_t^2] &= E[\epsilon_t^4] - E[\epsilon_t^2]^2 \\ &= \frac{3\omega^2(1+\alpha_1)}{(1-3\alpha_1^2)(1-\alpha_1)} - \frac{\omega^2}{(1-\alpha_1)^2} \\ &= \frac{3\omega^2(1+\alpha_1)(1-\alpha_1)^2}{(1-3\alpha_1^2)(1-\alpha_1)(1-\alpha_1)^2} - \frac{\omega^2(1-3\alpha_1^2)(1-\alpha_1)}{(1-3\alpha_1^2)(1-\alpha_1)(1-\alpha_1)^2} \\ &= \frac{3\omega^2(1+\alpha_1)(1-\alpha_1)^2 - \omega^2(1-3\alpha_1^2)(1-\alpha_1)}{(1-3\alpha_1^2)(1-\alpha_1)(1-\alpha_1)^2} \\ &= \frac{3\omega^2(1+\alpha_1)(1-\alpha_1) - \omega^2(1-3\alpha_1^2)}{(1-3\alpha_1^2)(1-\alpha_1)^2} \\ &= \frac{3\omega^2(1-\alpha_1^2) - \omega^2(1-3\alpha_1^2)}{(1-3\alpha_1^2)(1-\alpha_1)^2} \\ &= \frac{3\omega^2(1-\alpha_1^2) - 3\omega^2(\frac{1}{3} - \alpha_1^2)}{(1-3\alpha_1^2)(1-\alpha_1)^2} \\ &= \frac{3\omega^2[(1-\alpha_1^2) - (\frac{1}{3} - \alpha_1^2)]}{(1-3\alpha_1^2)(1-\alpha_1)^2} \\ &= \frac{2\omega^2}{(1-3\alpha_1^2)(1-\alpha_1)^2} \\ &= \left(\frac{\omega}{1-\alpha_1} \right)^2 \frac{2}{(1-3\alpha_1^2)} \\ &= \frac{2\bar{\sigma}^4}{(1-3\alpha_1^2)} \end{aligned} \tag{7.136}$$

The variance of the squared returns depends on the unconditional level of the variance, $\bar{\sigma}^2$, and the innovation term (α_1) squared.

7.B Kurtosis of a GARCH(1,1)

First, note that $E[\sigma_t^2 - \epsilon_t^2] = 0$, so that $V[\sigma_t^2 - \epsilon_t^2] = E[(\sigma_t^2 - \epsilon_t^2)^2]$. This term can be expanded to $E[\epsilon_t^4] - 2E[\epsilon_t^2\sigma_t^2] + E[\sigma_t^4]$ which can be shown to be $2E[\sigma_t^4]$ since

$$\begin{aligned} E[\epsilon_t^4] &= E[E_{t-1}[e_t^4\sigma_t^4]] \\ &= E[E_{t-1}[e_t^4]\sigma_t^4] \\ &= E[3\sigma_t^4] \\ &= 3E[\sigma_t^4] \end{aligned} \quad (7.137)$$

and

$$\begin{aligned} E[\epsilon_t^2\sigma_t^2] &= E[E_{t-1}[e_t^2\sigma_t^2]\sigma_t^2] \\ &= E[\sigma_t^2\sigma_t^2] \\ &= E[\sigma_t^4] \end{aligned} \quad (7.138)$$

so

$$\begin{aligned} E[\epsilon_t^4] - 2E[\epsilon_t^2\sigma_t^2] + E[\sigma_t^4] &= 3E[\sigma_t^4] - 2E[\sigma_t^4] + E[\sigma_t^4] \\ &= 2E[\sigma_t^4] \end{aligned} \quad (7.139)$$

The only remaining step is to complete the tedious derivation of the expectation of this fourth power,

$$\begin{aligned} E[\sigma_t^4] &= E[(\sigma_t^2)^2] \\ &= E[(\omega + \alpha_1\epsilon_{t-1}^2 + \beta_1\sigma_{t-1}^2)^2] \\ &= E[\omega^2 + 2\omega\alpha_1\epsilon_{t-1}^2 + 2\omega\beta_1\sigma_{t-1}^2 + 2\alpha_1\beta_1\epsilon_{t-1}^2\sigma_{t-1}^2 + \alpha_1^2\epsilon_{t-1}^4 + \beta_1^2\sigma_{t-1}^4] \\ &= \omega^2 + 2\omega\alpha_1E[\epsilon_{t-1}^2] + 2\omega\beta_1E[\sigma_{t-1}^2] + 2\alpha_1\beta_1E[\epsilon_{t-1}^2\sigma_{t-1}^2] + \alpha_1^2E[\epsilon_{t-1}^4] + \beta_1^2E[\sigma_{t-1}^4] \end{aligned} \quad (7.140)$$

Noting that

- $E[\epsilon_{t-1}^2] = E[E_{t-2}[\epsilon_{t-1}^2]] = E[E_{t-2}[e_{t-1}^2\sigma_{t-1}^2]] = E[\sigma_{t-1}^2E_{t-2}[e_{t-1}^2]] = E[\sigma_{t-1}^2] = \bar{\sigma}^2$
- $E[\epsilon_{t-1}^2\sigma_{t-1}^2] = E[E_{t-2}[\epsilon_{t-1}^2]\sigma_{t-1}^2] = E[E_{t-2}[e_{t-1}^2\sigma_{t-1}^2]\sigma_{t-1}^2] = E[E_{t-2}[e_{t-1}^2]\sigma_{t-1}^2\sigma_{t-1}^2] = E[\sigma_{t-1}^4]$
- $E[\epsilon_{t-1}^4] = E[E_{t-2}[\epsilon_{t-1}^4]] = E[E_{t-2}[e_{t-1}^4\sigma_{t-1}^4]] = 3E[\sigma_{t-1}^4]$

the final expression for $E[\sigma_t^4]$ can be arrived at

$$\begin{aligned} E[\sigma_t^4] &= \omega^2 + 2\omega\alpha_1E[\epsilon_{t-1}^2] + 2\omega\beta_1E[\sigma_{t-1}^2] + 2\alpha_1\beta_1E[\epsilon_{t-1}^2\sigma_{t-1}^2] + \alpha_1^2E[\epsilon_{t-1}^4] + \beta_1^2E[\sigma_{t-1}^4] \\ &= \omega^2 + 2\omega\alpha_1\bar{\sigma}^2 + 2\omega\beta_1\bar{\sigma}^2 + 2\alpha_1\beta_1E[\sigma_{t-1}^4] + 3\alpha_1^2E[\sigma_{t-1}^4] + \beta_1^2E[\sigma_{t-1}^4]. \end{aligned} \quad (7.141)$$

$E[\sigma_t^4]$ can be solved for (replacing $E[\sigma_t^4]$ with μ_4),

$$\begin{aligned}\mu_4 &= \omega^2 + 2\omega\alpha_1\bar{\sigma}^2 + 2\omega\beta_1\bar{\sigma}^2 + 2\alpha_1\beta_1\mu_4 + 3\alpha_1^2\mu_4 + \beta_1^2\mu_4 \quad (7.142) \\ \mu_4 - 2\alpha_1\beta_1\mu_4 - 3\alpha_1^2\mu_4 - \beta_1^2\mu_4 &= \omega^2 + 2\omega\alpha_1\bar{\sigma}^2 + 2\omega\beta_1\bar{\sigma}^2 \\ \mu_4(1 - 2\alpha_1\beta_1 - 3\alpha_1^2 - \beta_1^2) &= \omega^2 + 2\omega\alpha_1\bar{\sigma}^2 + 2\omega\beta_1\bar{\sigma}^2 \\ \mu_4 &= \frac{\omega^2 + 2\omega\alpha_1\bar{\sigma}^2 + 2\omega\beta_1\bar{\sigma}^2}{1 - 2\alpha_1\beta_1 - 3\alpha_1^2 - \beta_1^2}\end{aligned}$$

finally substituting $\bar{\sigma}^2 = \omega/(1 - \alpha_1 - \beta_1)$ and returning to the original derivation,

$$E[\epsilon_t^4] = \frac{3(1 + \alpha_1 + \beta_1)}{(1 - \alpha_1 - \beta_1)(1 - 2\alpha_1\beta_1 - 3\alpha_1^2 - \beta_1^2)}, \quad (7.143)$$

and the kurtosis, $\kappa = \frac{E[\epsilon_t^4]}{E[\epsilon_t^2]^2} = \frac{\mu_4}{\bar{\sigma}^2}$, which simplifies to

$$\kappa = \frac{3(1 + \alpha_1 + \beta_1)(1 - \alpha_1 - \beta_1)}{1 - 2\alpha_1\beta_1 - 3\alpha_1^2 - \beta_1^2} > 3. \quad (7.144)$$

Short Problems

Problem 7.1. What is Realized Variance and why is it useful?

Problem 7.2. Suppose $r_t = \sigma_t \epsilon_t$ where $\sigma_t^2 = \omega + \alpha r_{t-1}^2 + \beta \sigma_{t-1}^2$, and $\epsilon_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$. What conditions are required on the parameters ω , α , and β for r_t to be covariance stationary?

Problem 7.3. What is Realized Variance?

Problem 7.4. Discuss the properties of the generalized forecast error from a correctly specified volatility model.

Problem 7.5. Outline the steps in the Mincer-Zarnowitz framework to objectively evaluate a sequence of variance forecasts $\{\hat{\sigma}_{t+1|t}^2\}$.

Exercises

Exercise 7.1. Suppose we model log-prices at time t , written p_t , as an ARCH(1) process

$$p_t | \mathcal{F}_{t-1} \sim N(p_{t-1}, \sigma_t^2),$$

where \mathcal{F}_t denotes the information up to and including time t and

$$\sigma_t^2 = \alpha + \beta (p_{t-1} - p_{t-2})^2.$$

1. Is p_t a martingale?

2. What is

$$E[\sigma_t^2]?$$

3. Calculate

$$\text{Cov} [(p_t - p_{t-1})^2, (p_{t-s} - p_{t-1-s})^2]$$

for $s > 0$.

4. Comment on the importance of this result from a practical perspective.

5. How do you use a likelihood function to estimate an ARCH model?

6. How can the ARCH(1) model be generalized better capture the variance dynamics of asset prices?

7. In the ARCH(1) case, what can you say about the properties of

$$p_{t+s} | \mathcal{F}_{t-1},$$

for $s > 0$, i.e., the multi-step ahead forecast of prices?

8. Why are Bollerslev-Wooldridge standard errors important when testing coefficients in ARCH models?

Exercise 7.2. Derive explicit relationships between the parameters of an APARCH(1,1,1),

$$\begin{aligned} r_t &= \mu_t + \epsilon_t \\ \sigma_t^\delta &= \omega + \alpha_1 (|\epsilon_{t-1}| + \gamma_1 \epsilon_{t-1})^\delta + \beta_1 \sigma_{t-1}^\delta \\ \epsilon_t &= \sigma_t e_t \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1), \end{aligned}$$

and:

1. ARCH(1)
2. GARCH(1,1)
3. AVGARCH(1,1)
4. TARCH(1,1,1)
5. GJR-GARCH(1,1,1)

Exercise 7.3. Consider the following GJR-GARCH process,

$$\begin{aligned} r_t &= \rho r_{t-1} + \epsilon_t \\ \epsilon_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha \epsilon_{t-1}^2 + \gamma \epsilon_{t-1}^2 I_{[\epsilon_{t-1} < 0]} + \beta \sigma_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

where $E_t[\cdot] = E[\cdot | \mathcal{F}_t]$ is the time t conditional expectation and $V_t[\cdot] = V[\cdot | \mathcal{F}_t]$ is the time t conditional variance.

1. What conditions are necessary for this process to be covariance stationary?

Assume these conditions hold in the remaining questions. *Note:* If you cannot answer one or more of these questions for an arbitrary γ , you can assume that $\gamma = 0$ and receive partial credit.

2. What is $E[r_{t+1}]$?
3. What is $E_t[r_{t+1}]$?
4. What is $V[r_{t+1}]$?
5. What is $V_t[r_{t+1}]$?
6. What is $V_t[r_{t+2}]$?

Exercise 7.4. Let r_t follow a GARCH process

$$\begin{aligned} r_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha r_{t-1}^2 + \beta \sigma_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

1. What are the values of the following quantities?
 - (a) $E[r_{t+1}]$
 - (b) $E_t[r_{t+1}]$
 - (c) $V[r_{t+1}]$
 - (d) $V_t[r_{t+1}]$
 - (e) $\rho_1 = \text{Corr}[r_t, r_{t-1}]$
2. What is $E[(r_t^2 - \bar{\sigma}^2)(r_{t-1}^2 - \bar{\sigma}^2)]$ where $\bar{\sigma} = E[\sigma_t^2]$. Hint: Consider the relationship to ARMA models.
3. Describe the h -step ahead forecast from this model.

Exercise 7.5. Let r_t follow an ARCH process

$$\begin{aligned} r_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha_1 r_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

1. What are the values of the following quantities?
 - (a) $E[r_{t+1}]$
 - (b) $E_t[r_{t+1}]$
 - (c) $V[r_{t+1}]$

- (d) $V_t[r_{t+1}]$
- (e) $\rho_1 = \text{Corr}[r_t, r_{t-1}]$
- 2. What is $E[(r_t^2 - \bar{\sigma}^2)(r_{t-1}^2 - \bar{\sigma}^2)]$ where $\bar{\sigma} = E[\sigma_t^2]$. Hint: Think about the AR duality.
- 3. Describe the h -step ahead forecast from this model.

Exercise 7.6. Consider an EGARCH(1,1,1) model:

$$\ln \sigma_t^2 = \omega + \alpha_1 \left(|e_{t-1}| - \sqrt{\frac{2}{\pi}} \right) + \gamma_1 e_{t-1} + \beta_1 \ln \sigma_{t-1}^2$$

where $e_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$.

1. What are the required conditions on the model parameters for this process to be covariance stationary?
2. What is the one-step-ahead forecast of σ_t^2 , $E_t [\sigma_{t+1}^2]$?
3. What is the most you can say about the two-step-ahead forecast of σ_t^2 ($E_t [\sigma_{t+2}^2]$)?

Exercise 7.7. Answer the following questions:

1. Describe three fundamentally different procedures to estimate the volatility over some interval. What the strengths and weaknesses of each?
2. Why is Realized Variance useful when evaluating a volatility model?
3. What considerations are important when computing Realized Variance?
4. Why does the Black-Scholes implied volatility vary across strikes?

Exercise 7.8. Consider a general volatility specification for an asset return r_t :

$$\begin{aligned} r_t | \mathcal{F}_{t-1} &\sim N(0, \sigma_t^2) \\ \text{and let } e_t &\equiv \frac{r_t}{\sigma_t} \\ \text{so } e_t | \mathcal{F}_{t-1} &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

1. Find the conditional kurtosis of the returns:

$$\text{Kurt}_{t-1}[r_t] \equiv \frac{E_{t-1} [(r_t - E_{t-1}[r_t])^4]}{(V_{t-1}[r_t])^2}$$

2. Show that if $V[\sigma_t^2] > 0$, then the *unconditional* kurtosis of the returns,

$$\text{Kurt}[r_t] \equiv \frac{E[(r_t - E[r_t])^4]}{(V[r_t])^2}$$

is greater than 3.

3. Find the conditional skewness of the returns:

$$\text{Skew}_{t-1}[r_t] \equiv \frac{E_{t-1}[(r_t - E_{t-1}[r_t])^3]}{(V_{t-1}[r_t])^{3/2}}$$

4. Find the *unconditional* skewness of the returns:

$$\text{Skew}[r_t] \equiv \frac{E[(r_t - E[r_t])^3]}{(V[r_t])^{3/2}}$$

Exercise 7.9. Answer the following questions:

1. Describe three fundamentally different procedures to estimate the volatility over some interval. What are the strengths and weaknesses of each?
2. Why does the Black-Scholes implied volatility vary across strikes?
3. Consider the following GJR-GARCH process,

$$\begin{aligned} r_t &= \mu + \rho r_{t-1} + \epsilon_t \\ \epsilon_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha \epsilon_{t-1}^2 + \gamma \epsilon_{t-1}^2 I_{[\epsilon_{t-1} < 0]} + \beta \sigma_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

where $E_t[\cdot] = E[\cdot | \mathcal{F}_t]$ is the time t conditional expectation and $V_t[\cdot] = V[\cdot | \mathcal{F}_t]$ is the time t conditional variance.

- (a) What conditions are necessary for this process to be covariance stationary?

Assume these conditions hold in the remaining questions.

- (b) What is $E[r_{t+1}]$?
- (c) What is $E_t[r_{t+1}]$?
- (d) What is $E_t[r_{t+2}]$?
- (e) What is $V[r_{t+1}]$?
- (f) What is $V_t[r_{t+1}]$?
- (g) What is $V_t[r_{t+2}]$?

Exercise 7.10. Answer the following questions about variance estimation.

1. What is Realized Variance?
2. How is Realized Variance estimated?
3. Describe two models which are appropriate for modeling Realized Variance.
4. What is an Exponential Weighted Moving Average (EWMA)?

5. Suppose an ARCH model for the conditional variance of daily returns was fit

$$\begin{aligned} r_{t+1} &= \mu + \sigma_{t+1} e_{t+1} \\ \sigma_{t+1}^2 &= \omega + \alpha_1 \epsilon_t^2 + \alpha_2 \epsilon_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

What are the forecasts for $t + 1$, $t + 2$ and $t + 3$ given the current (time t) information set?

6. Suppose an EWMA was used instead for the model of conditional variance with smoothing parameter = .94. What are the forecasts for $t + 1$, $t + 2$ and $t + 3$ given the current (time t) information set?
7. Compare the ARCH(2) and EWMA forecasts when the forecast horizon is large (e.g., $E_t [\sigma_{t+h}^2]$ for large h).
8. What is VIX?

Exercise 7.11. Suppose $\{y_t\}$ is covariance stationary and can be described by the following process:

$$\begin{aligned} y_t &= \phi_1 y_{t-1} + \theta_1 \epsilon_{t-1} + \epsilon_t \\ \epsilon_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha \epsilon_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

what are the values of the following quantities:

1. $E_t [y_{t+1}]$
2. $E_t [y_{t+2}]$
3. $\lim_{h \rightarrow \infty} E_t [y_{t+h}]$
4. $V_t [\epsilon_{t+1}]$
5. $V_t [y_{t+1}]$
6. $V_t [y_{t+2}]$
7. $\lim_{h \rightarrow \infty} V_t [\epsilon_{t+h}]$

Exercise 7.12. Answer the following questions:

Suppose $\{y_t\}$ is covariance stationary and can be described by the following process:

$$\begin{aligned} y_t &= \phi_0 + \theta_1 \epsilon_{t-1} + \theta_2 \epsilon_{t-2} + \epsilon_t \\ \epsilon_t &= \sigma_t e_t \\ \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \beta \sigma_{t-1}^2 \\ e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1) \end{aligned}$$

what are the values of the following quantities:

1. $E_t[y_{t+1}]$
2. $E_t[y_{t+2}]$
3. $\lim_{h \rightarrow \infty} E_t[y_{t+h}]$
4. $V_t[\epsilon_{t+1}]$
5. $V_t[y_{t+2}]$
6. $\lim_{h \rightarrow \infty} V_t[\epsilon_{t+h}]$

Exercise 7.13. Consider the AR(2)-ARCH(2) model

$$\begin{aligned}
 y_t &= \phi_0 + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \epsilon_t \\
 \epsilon_t &= \sigma_t e_t \\
 \sigma_t^2 &= \omega + \alpha_1 \epsilon_{t-1}^2 + \alpha_2 \epsilon_{t-2}^2 \\
 e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1)
 \end{aligned}$$

1. What conditions are required for ϕ_0 , ϕ_1 and, ϕ_2 for the model to be covariance stationary?
2. What conditions are required for ω , α_1 , and α_2 for the model to be covariance stationary?
3. Show that $\{\epsilon_t\}$ is a white noise process.
4. Are ϵ_t and ϵ_{t-s} independent for $s \neq 0$?
5. What are the values of the following quantities:
 - (a) $E[y_t]$
 - (b) $E_t[y_{t+1}]$
 - (c) $E_t[y_{t+2}]$
 - (d) $V_t[y_{t+1}]$
 - (e) $V_t[y_{t+2}]$

Exercise 7.14. Suppose $\{y_t\}$ is covariance stationary and can be described by the following process:

$$\begin{aligned}
 y_t &= \phi_1 y_{t-1} + \theta_1 \epsilon_{t-1} + \epsilon_t \\
 \epsilon_t &= \sigma_t e_t \\
 \sigma_t^2 &= \omega + \alpha \epsilon_{t-1}^2 \\
 e_t &\stackrel{\text{i.i.d.}}{\sim} N(0, 1)
 \end{aligned}$$

1. What are the values of the following quantities:
 - (a) $E_t[y_{t+1}]$
 - (b) $E_t[y_{t+2}]$

- (c) $\lim_{h \rightarrow \infty} E_t [y_{t+h}]$
- (d) $V_t [\epsilon_{t+1}]$
- (e) $V_t [y_{t+1}]$
- (f) $V_t [y_{t+2}]$
- (g) $V [y_{t+1}]$

Exercise 7.15. Consider the MA(2)-GARCH(1,1) model

$$y_t = \phi_0 + \theta_1 \epsilon_{t-1} + \phi_2 \epsilon_{t-2} + \epsilon_t$$

$$\epsilon_t = \sigma_t e_t$$

$$\sigma_t^2 = \omega + \alpha_1 \epsilon_{t-1}^2 + \beta_1 \sigma_{t-1}^2$$

$$e_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$$

1. What conditions are required for ϕ_0 , θ_1 , and θ_2 for the model to be covariance stationary?
2. What conditions are required for ω, α_1 , and β_1 for the model to be covariance stationary?
3. Show that $\{\epsilon_t\}$ is a white noise process.
4. Are ϵ_t and ϵ_{t-1} independent?
5. What are the values of the following quantities:

- (a) $E[y_t]$
- (b) $E_t [y_{t+1}]$
- (c) $E_t [y_{t+2}]$
- (d) $\lim_{h \rightarrow \infty} E_t [y_{t+h}]$
- (e) $V_t [y_{t+1}]$
- (f) $V_t [y_{t+2}]$

Exercise 7.16. Suppose $\{y_t\}$ is covariance stationary and can be described by the following process:

$$y_t = \phi_0 + \phi_1 y_{t-1} + \theta_1 \epsilon_{t-1} + \epsilon_t$$

$$\epsilon_t = \sigma_t e_t$$

$$\sigma_t^2 = \omega + \alpha \epsilon_{t-1}^2$$

$$e_t \stackrel{\text{i.i.d.}}{\sim} N(0, 1)$$

1. What are the values of the following quantities:

- (a) $E[y_{t+1}]$
- (b) $E_t [y_{t+1}]$
- (c) $E_t [y_{t+2}]$

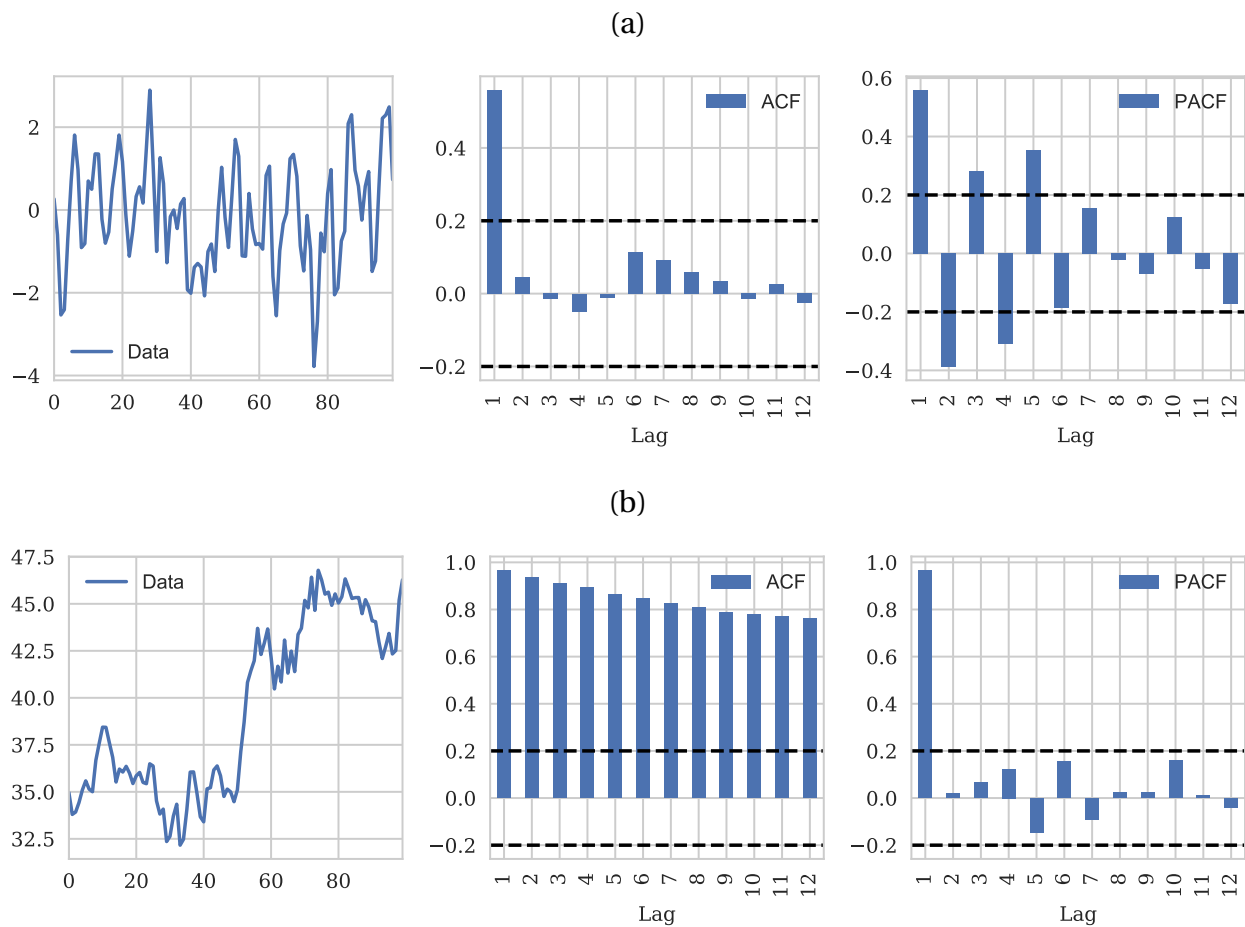


Figure 7.15: Plots for question 7.16.

- (d) $\lim_{h \rightarrow \infty} E_t[y_{t+h}]$
- (e) $V_t[\epsilon_{t+1}]$
- (f) $V_t[y_{t+1}]$
- (g) $V_t[y_{t+2}]$
- (h) $V[y_{t+1}]$

2. Justify a reasonable model for each of these time series in Figure 7.15 using information in the autocorrelation and partial autocorrelation plots. In each set of plots, the leftmost panel shows that data ($T = 100$). The middle panel shows the sample autocorrelation with 95% confidence bands. The right panel shows the sample partial autocorrelation for the data with 95% confidence bands.

- (a) Panel (a)
- (b) Panel (b)

